

Document stores

Taxonomy of NoSQL

- Key-value



- Graph database




- Document-oriented



- Column family



Typical NoSQL Architecture



Document stores

- Flexible schema
- JSON/BSON documents
 - Embedded documents
 - Referenced documents
- CouchDB: <http://couchdb.apache.org/>
- MongoDB: <http://www.mongodb.org/>

We'll use the mongo shell for class, but if you want to use a GUI to interact with MongoDB, you may want to look into Robo 3T (previously robomongo) <https://robomongo.org/>

What is MongoDB?

- Developed by 10gen
 - Founded in 2007
- A document-oriented, NoSQL database
 - Hash-based, *schema-less database*
 - No Data Definition Language
 - In practice, this means you can store hashes with any keys and values that you choose
 - Keys are a basic data type but in reality stored as strings
 - Document Identifiers (`_id`) will be created for each document, field name reserved by system
 - Application tracks the schema and mapping
 - Uses BSON format
 - Based on JSON – B stands for Binary
- Written in C++
- Supports APIs (drivers) in many computer languages

MongoDB Features

- Dynamic schema
- Document-Oriented storage
- Full Index Support
- Replication & High Availability
- Auto-Sharding
 - Built-in horizontal scaling via automated range-based partitioning of data
- Querying
- Fast In-Place Updates
- Map/Reduce functionality


Agile

Scalable

MongoDB: CAP approach

Focus on Consistency and Partition tolerance


- **Consistency**
 - all replicas contain the same version of the data
- **Availability**
 - system remains operational on failing nodes
- **Partition tolerance**
 - multiple entry points
 - system remains operational on system split



CAP Theorem:
satisfying all three at the same time is
impossible

MongoDB: Hierarchical Objects

- A MongoDB instance may have zero or more 'databases'
- A database may have zero or more 'collections'.
- A collection may have zero or more 'documents'.
- A document may have one or more 'fields'.
- MongoDB 'Indexes' function much like their RDBMS counterparts.



MongoDB

RDBMS		MongoDB
Database	➡	Database
Table, View	➡	Collection
Row	➡	Document (JSON, BSON)
Column	➡	Field
Index	➡	Index
Join	➡	Embedded Document
Foreign Key	➡	Reference
Partition	➡	Shard

Mongo hits a sweet spot between the powerful queryability of a relational database and the distributed nature of other databases

JSON format

- Data is in name / valuepairs
- A name/value pair consists of a field name followed by a colon, followed by a value:
 - Example: "name": "R2-D2"
- Data is separated by commas
 - Example: "name": "R2-D2", race : "Droid"
- Curly braces hold objects
 - Example: {"name": "R2-D2", race : "Droid", affiliation: "rebels"}
- An array is stored in brackets []
 - Example [{"name": "R2-D2", race : "Droid", affiliation: "rebels"}, {"name": "Yoda", affiliation: "rebels"}]

Document store

RDBMS		MongoDB
Database	➡	Database
Table, View	➡	Collection
Row	➡	Document (JSON, BSON)
Column	➡	Field
Index	➡	Index
Join	➡	Embedded Document
Foreign Key	➡	Reference
Partition	➡	Shard

```
> db.user.findOne({age:39})
{
  "_id" : ObjectId("5114e0bd42..."),
  "first" : "John",
  "last" : "Doe",
  "age" : 39,
  "interests" : [
    "Reading",
    "Mountain Biking"
  ],
  "favorites": {
    "color": "Blue",
    "sport": "Soccer"
  }
}
```

CRUD

Create

```
db.collection.insert( <document> )
```

```
db.collection.save( <document> )
```

```
db.collection.update( <query>, <update>, { upsert: true } )
```

Read

```
db.collection.find( <query>, <projection> )
```

```
db.collection.findOne( <query>, <projection> )
```

Update

```
db.collection.update( <query>, <update>, <options> )
```

Delete

```
db.collection.remove( <query>, <justOne> )
```

CRUD example

```
> db.user.insert({  
  first: "John",  
  last : "Doe",  
  age: 39  
})
```

```
> db.user.find ()  
{  
  "_id" : ObjectId("51..."),  
  "first" : "John",  
  "last" : "Doe",  
  "age" : 39  
}
```

```
> db.user.update(  
  {"_id" : ObjectId("51...")},  
  {  
    $set: {  
      age: 40,  
      salary: 7000}  
    }  
  )
```

```
> db.user.remove({  
  "first": /^J/  
})
```

Let's get started – command-line fun

```
> mongo moderndb
```

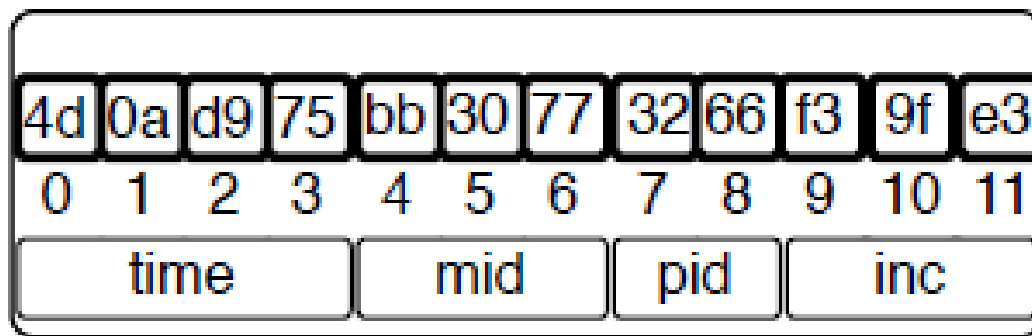
```
> db.towns.insert({  
  name: "New York", population:  
  22200000,  
  lastCensus: ISODate("2016-07-  
  01"),  
  famousFor: [ "the MOMA",  
    "food", "Derek Jeter" ], mayor  
  : {  
    name : "Bill de Blasio",  
    party : "D"  
  }  
})
```

```
> show collections
```

```
> db.towns.find()
```

ObjectId

- `_id` field of type ObjectId.
- Akin to SERIAL incrementing a numeric primary key in PostgreSQL.
- The ObjectId is always 12 bytes, composed of a timestamp, client machine ID, client process ID, and a 3-byte incremented counter.



Javascript

- Native tongue of MongoDB
- Ask for help
 - `db.help()`
 - `db.towns.help()`
- Object and functions
 - `typeof db`
 - `typeof db.towns`
 - `typeof db.towns.insert`
- Get source code of a function (call it without parameters)
 - `db.towns.insert`

Let's insert data using our own function

```
function insertCity(name, population, lastCensus, famousFor,
mayorInfo) {
  db.towns.insert({ name: name,
    population: population, lastCensus: ISODate(lastCensus),
    famousFor: famousFor,
    mayor : mayorInfo });
}
```

Now we can call it

```
> insertCity("Punxsutawney", 6200, '2016-01-31', ["Punxsutawney
Phil"], { name : "Richard Alexander" }
)
```

```
> insertCity("Portland", 582000, '2016-09-20',
["beer", "food", "Portlandia"], { name : "Ted Wheeler", party :
"D" })
```

Querying data

```
db.towns.find({ "_id" : ObjectId( "59094288afbc9350ada6b807" ) })
```

```
db.towns.find({ _id : ObjectId( "59094288afbc9350ada6b807" ) }, {  
name : 1 })
```

```
db.towns.find({ _id : ObjectId( "59094288afbc9350ada6b807" ) }, {  
name : 0 })
```

Perl-compatible regular expression (PCRE)

```
db.towns.find(  
{ name : /^P/, population : { $lt : 10000 } },  
{ _id: 0, name : 1, population : 1 }  
)
```

Can construct operations as you would objects

```
> var population_range = {$lt: 1000000, $gt: 10000}  
> db.towns.find( { name : /^P/, population : population_range }, {  
name: 1 })
```

```
> db.towns.find(  
{ lastCensus : { $gte : ISODate( '2016-06-01' ) } },  
{ _id : 0, name: 1 })
```

Querying nested array data

Matching exact values

```
> db.towns.find(  
  { famousFor : 'food' },  
  { _id : 0, name : 1, famousFor : 1 }  
)
```

Matching partial values:

```
> db.towns.find(  
  { famousFor : /moma/ },  
  { _id : 0, name : 1, famousFor : 1 })
```

Query by all matching values:

```
> db.towns.find(  
  { famousFor : { $all : [ 'food', 'beer' ] } },  
  { _id : 0, name : 1, famousFor : 1 }  
)
```

Or the lack of matching values:

```
> db.towns.find(  
  { famousFor : { $nin : [ 'food', 'beer' ] } },  
  { _id : 0, name : 1, famousFor : 1 }  
)
```

Querying nested documents

Find towns with mayors from the Democratic party

```
> db.towns.find(  
  { 'mayor.party' : 'D' },  
  { _id : 0, name : 1, mayor : 1 }
```

Find towns with mayors who don't have a party

```
> db.towns.find(  
  { 'mayor.party' : { $exists : false } },  
  { _id : 0, name : 1, mayor : 1 }  
)
```

New collection for countries

```
> db.countries.insert({
  _id : "us",
  name : "United States",
  exports : {
    foods : [
      { name : "bacon", tasty : true },
      { name : "burgers" }
    ]
  }
})
> db.countries.insert({
  _id : "ca",
  name : "Canada",
  exports : {
    foods : [
      { name : "bacon", tasty : false },
      { name : "syrup", tasty : true }
    ]
  }
})
> db.countries.insert({
  _id : "mx",
  name : "Mexico",
  exports : {
    foods : [{
      name : "salsa", tasty : true, condiment : true
    }]
  }
})
```

elemMatch

Find a country that not only exports bacon but exports tasty bacon

```
> db.countries.find(  
  { 'exports.foods.name' : 'bacon', 'exports.foods.tasty' : true },  
  { _id : 0, name : 1 }  
)
```

Canada?

ElemMatch to the rescue:

```
> db.countries.find(  
  {  
    'exports.foods' : {  
      $elemMatch : {  
        name : 'bacon',  
        tasty : true  
      }  
    },  
    { _id : 0, name : 1 }  
  })
```

```
> db.countries.find(  
  {  
    'exports.foods' : {  
      $elemMatch : {  
        tasty : true,  
        condiment : { $exists : true }  
      }  
    },  
    { _id : 0, name : 1 }  
  })
```

Boolean operators

```
> db.countries.find(  
  { _id: "mx", name: "United States" },  
  { _id: 1 }  
)
```

```
db.countries.find(  
  {  
    $or: [  
      { _id: "mx" },  
      { name: "United States" }  
    ]  
  },  
  { _id: 1 }  
)
```

Some mongodb commands

Command	Description
\$regex	Match by any PCRE-compliant regular expression string (or just use the // delimiters as shown earlier)
\$ne	Not equal to
\$lt	Less than
\$lte	Less than or equal to
\$gt	Greater than
\$gte	Greater than or equal to
\$exists	Check for the existence of a field
\$all	Match all elements in an array
\$in	Match any elements in an array
\$nin	Does not match any elements in an array
\$elemMatch	Match all fields in an array of nested documents
\$or	or
\$nor	Not or
\$size	Match array of given size
\$mod	Modulus
\$type	Match if field is a given datatype
\$not	Negate the given operator check

- Check the mongo documentation for a complete list
- Cheat sheet on ICON

For today..

- Select a town via a case-insensitive regular expression containing the word *new*.
- Find all towns whose names contain an *e* and are famous for food or beer.
- Submit the two queries to ICON.