

Indexing and aggregation

MongoDB



Import json file to MongoDB

<https://docs.mongodb.com/guides/server/import/>

If you didn't import the city_inspections last class:

- Download the file city_inspections.json from ICON
- Use the mongoimport utility:

```
mongoimport --db moderndb --collection city_inspections  
            --drop --file ~\downloads\city_inspections.json
```

Or if you enabled authentication:

```
mongoimport --db moderndb --collection city_inspections  
            --authenticationDatabase admin --username <user>  
            --password <password> --drop  
            --file ~\downloads\city_inspections.json
```

Execution stats

Let's get the execution statistics for a query by certificate_number in the city_inspections collection:

```
db.city_inspections.find({certificate_number:  
10003581}).explain("executionStats").executionStats
```

- You can also set the profiler to record all or long running queries. To disactivate the profiler (default)

```
db.setProfilingLevel(0)
```

- Recall that Mongo automatically creates an index by the _id

```
db.getCollectionNames().forEach(function(collection) {  
    print("Indexes for the " + collection + " collection:");  
    printjson(db[collection].getIndexes());  
});
```

Create an Index

Let's create an index on `certificate_number`

```
db.city_inspections.createIndex ({“certificate_number”: 1}, {unique: false})
```

Now let's see the improvement on executing the same query:

```
db.city_inspections.find({certificate_number:  
10003581}).explain("executionStats").executionStats
```

Query by id

```
db.city_inspections.find({id:"108-2015-  
UNIT"}).explain("executionStats").executionStats
```

Now let's create a hash index over id

```
db.city_inspections.createIndex ({“id”: “hashed”})
```

Create an index on zip code (from address document)

```
db.city_inspections.findOne()
```

Ascending order

```
db.city_inspections.createIndex ({“address.zip”: 1})
```

Descending order

```
db.city_inspections.createIndex ({“address.zip”: -1})
```

```
db.city_inspections.getIndexes()
```

Let's drop one:

```
db.city_inspections.dropIndex (“address.zip_1”)
```

More queries

```
db.city_inspections.find ({"certificate_number": {"$lt": 100000}}).sort(
{"address.zip": -1})
```

Single-purpose aggregators:

Count - number of documents in the result

```
db.city_inspections.find ({"certificate_number": {"$lt": 100000}}).count()
db.city_inspections.count ({"certificate_number": {"$lt": 100000}})
```

Distinct – collect the result set into an array of unique values

```
db.city_inspections.distinct("address.zip",
{"certificate_number": {"$gt": 100000}})
```

Aggregate – returns document according to the logic you provide

Aggregate (a pipeline-style logic)

<https://docs.mongodb.com/manual/aggregation/#aggregation-pipeline>

Stages (full list <https://docs.mongodb.com/manual/reference/operator/aggregation-pipeline/>):

\$match – filters

\$group – group by

\$sort – order by

\$project – select tags/documents to display in the results set

\$limit – limit the number of results

* hint option can be used to force the usage of the specified index

Count the number of passed inspections per city

```
db.city_inspectors.aggregate([
  {$match: { "result": {$eq: "No Violation Issued"} }},
  {$group: {
    _id: "$address.city",
    count: {$sum: 1}
  }}
])
```


Count the number of passed inspections per city order by count

```
db.city_inspections.aggregate([
  {$match: { "result": {$eq: "No Violation Issued"} }},
  {$group: {
    _id: "$address.city",
    count: {$sum: 1}
  }},
  {$sort: {"count": -1}}
])
```

Cities with over 200 passed inspections order by count

```
db.ci ty_i nspections.aggregate([
  {$match: { "result": {$eq: "No Vi ol at ion Issued"}}},
  {$group: {
    _i d: "$address. ci ty",
    count: {$sum: 1}
  }
},
  {$sort: {"count": - 1}},
  {$match: {"count": {$gt: 200}}}
])
```

Server side commands

- Pre-built Mongo commands execute in the server
- Some of them need to be executed under the admin database
- use `use moderndb`
- `db.listCommands()` –most of the commands execute on the server, not the client
- To have our own functions executed on the server (similar to stored procedures), add it to `collection.system.js`

```
db.system.js.save ({_id: "getLast", value: function(collection) {  
    return collection.find({}).sort({'_id': 1}).limit(1)[0];  
    }  
})  
use moderndb  
db.loadServerScripts()  
getLast(db.inventory).display
```

For today..

- <https://docs.mongodb.com/manual/tutorial/aggregation-zip-code-data-set/>
- Write a query to return Largest and Smallest Cities for these Midwest States: **Illinois**, **Indiana**, Iowa, and **Kansas**
- Submit your query to ICON