

Map Reduce

Background

- Google deals with very large amounts of data (petabytes)
 - need to process data fairly quickly
 - use very large numbers of commodity machines
 - Cheap nodes fail, especially if you have many
 - Mean time between failures for 1 node = 3 years
 - Mean time between failures for 1000 nodes = 1 day
 - Solution: Build fault-tolerance into the system
- Google developed an infrastructure consisting of
 - the Google distributed file system GFS
 - the MapReduce computational model
- MapReduce
 - functional programming model
 - Automatic parallelization & distribution
 - Fault tolerance
 - I/O scheduling
 - Monitoring & status updates
- Open source implementation Hadoop from Apache

Map Reduce

- Programming model for indexing and searching large data volumes over computer clusters
- Two Phases, Map and Reduce
 - Map
 - Extract sets of Key-Value pairs from underlying data
 - Potentially in Parallel on multiple machines
 - Reduce
 - Merge and sort sets of Key-Value pairs
 - Results may be useful for other searches

Google MapReduce

- Google's MapReduce is implemented as a C++ library.
- Operates on commodity hardware and standard networking.
- Input data, intermediate results, and final results are stored in GFS.
- A master scheduler process distributes map, reduce tasks to workers.
- Fault tolerance:
 - The master pings workers periodically.
 - Workers that do not respond are marked as failed.
 - Jobs assigned to failed workers are rerun.
 - Master failure aborts the computation.

Data type: key-value *records*

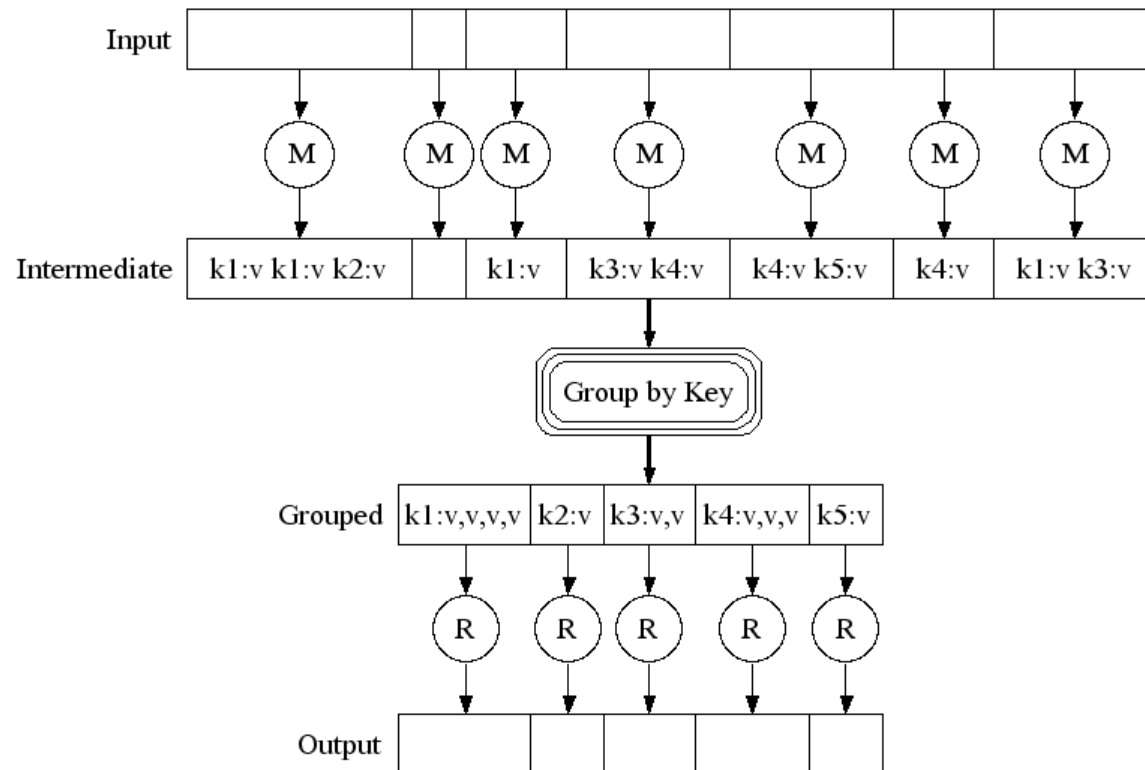
Map function:

$$(K_{in}, V_{in}) \rightarrow \text{list}(K_{inter}, V_{inter})$$

Reduce function:

$$(K_{inter}, \text{list}(V_{inter})) \rightarrow \text{list}(K_{out}, V_{out})$$

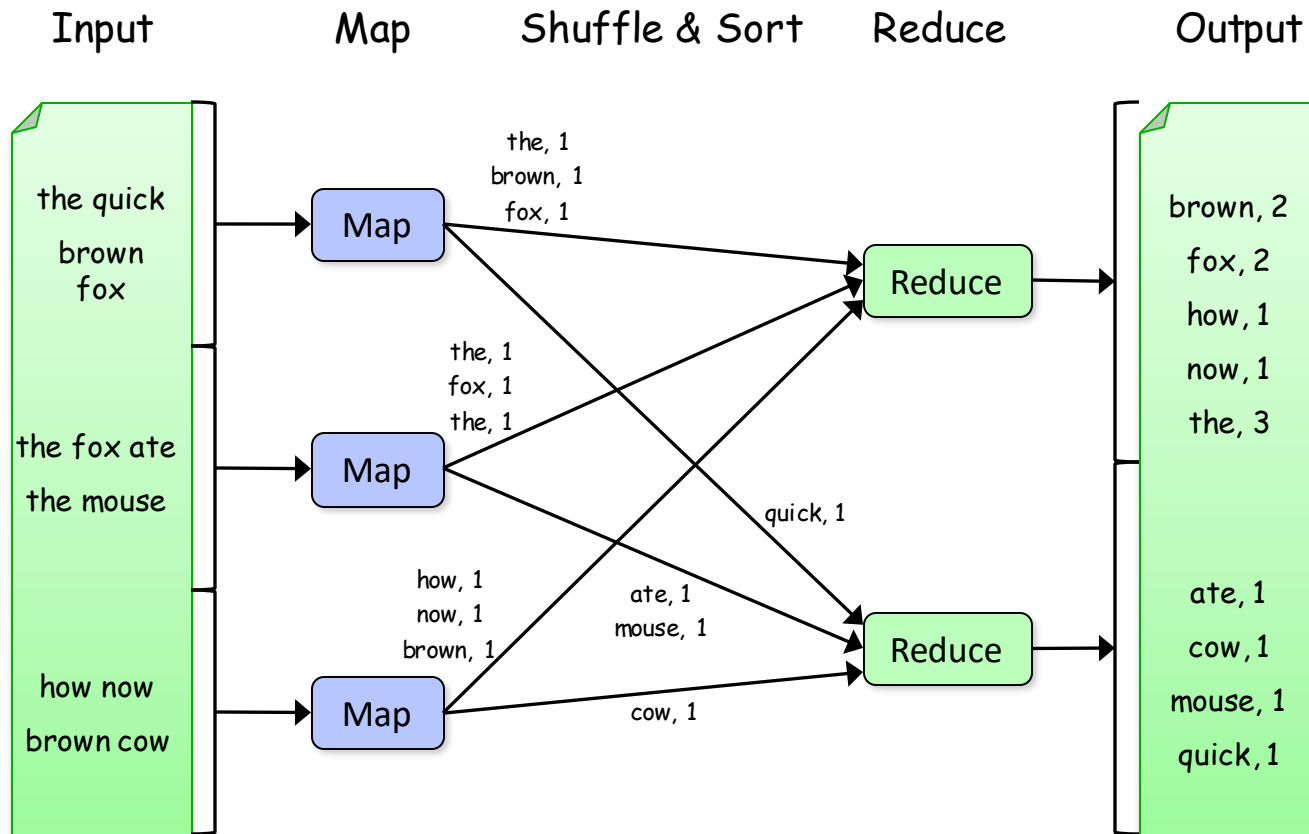
Execution



Example: Word Count

```
def mapper(line):  
    foreach word in line.split():  
        output(word, 1)  
  
def reducer(key, values):  
    output(key, sum(values))
```

Word Count Execution

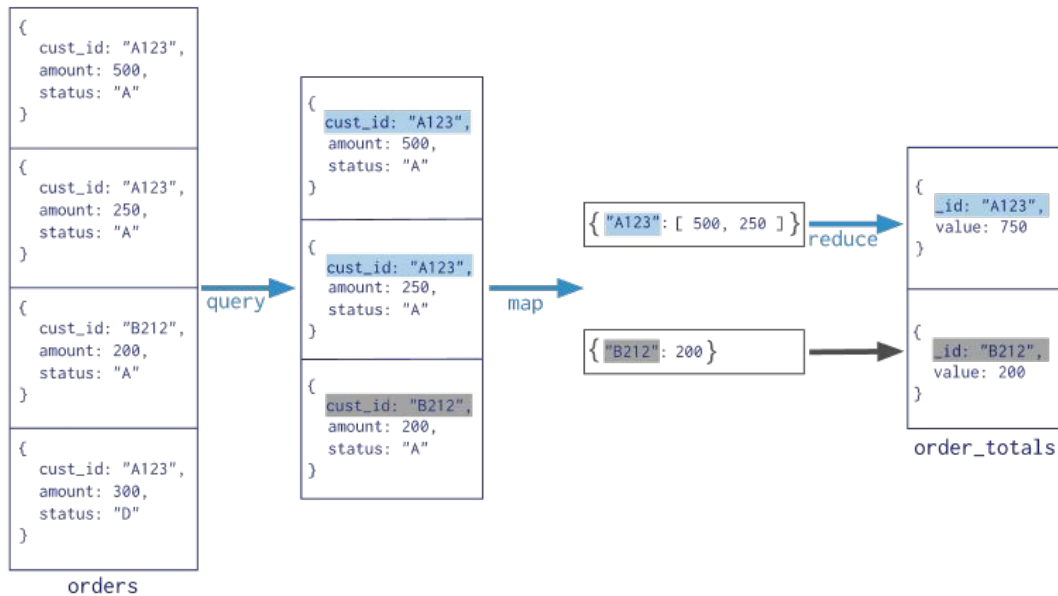


MapReduce Execution Details

- Single *master* controls job execution on multiple *slaves*
- Mappers preferentially placed on same node or same rack as their input block
 - Minimizes network usage
- Mappers save outputs to local disk before serving them to reducers
 - Allows recovery if a reducer crashes
 - Allows having more reducers than nodes

Map Reduce in MongoDB

Collection
↓
db.orders.mapReduce(
 map → function() { emit(this.cust_id, this.amount); },
 reduce → function(key, values) { return Array.sum(values) },
 query → {
 query: { status: "A" },
 output: "order_totals"
 }
)



Map/Reduce in Mongo

```
db.collection.mapReduce(  
    <mapfunction>,  
    <reducefunction>,  
    {  
        out: <collection>,  
        query: <>,  
        sort: <>,  
        limit: <number>,  
        finalize: <function>,  
        scope: <>,  
        jsMode: <boolean>,  
        verbose: <boolean>  
    }  
)
```

Example: Tickets

```
{  
  "id": 1,  
  "day": 20100123,  
  "checkout": 100  
}
```

```
{  
  "id": 2,  
  "day": 20100123,  
  "checkout": 42  
}
```

```
{  
  "id": 3,  
  "day": 20100123,  
  "checkout": 215  
}
```

```
{  
  "id": 4,  
  "day": 20100123,  
  "checkout": 73  
}
```

Sum(checkout)?

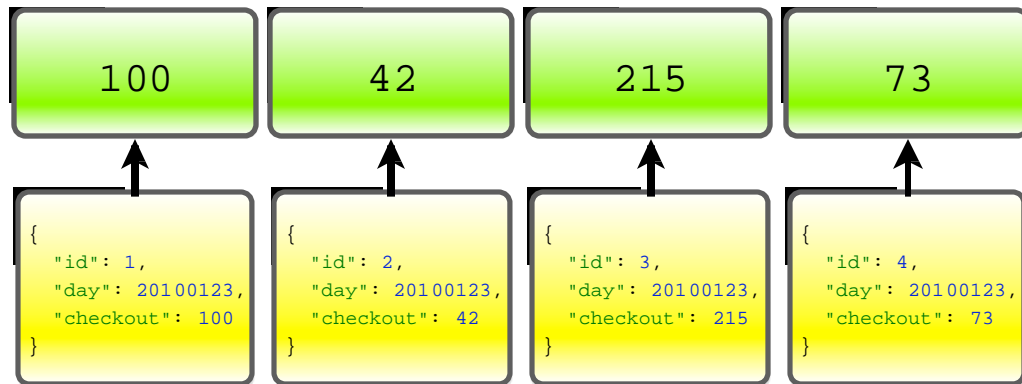
```
{  
  "id": 1,  
  "day": 20100123,  
  "checkout": 100  
}
```

```
{  
  "id": 2,  
  "day": 20100123,  
  "checkout": 42  
}
```


```
{  
  "id": 3,  
  "day": 20100123,  
  "checkout": 215  
}
```

```
{  
  "id": 4,  
  "day": 20100123,  
  "checkout": 73  
}
```


Map: emit(checkout)



Reduce: $\text{sum}(\text{checkouts})$



Reduce: $\text{sum}(\text{checkouts})$



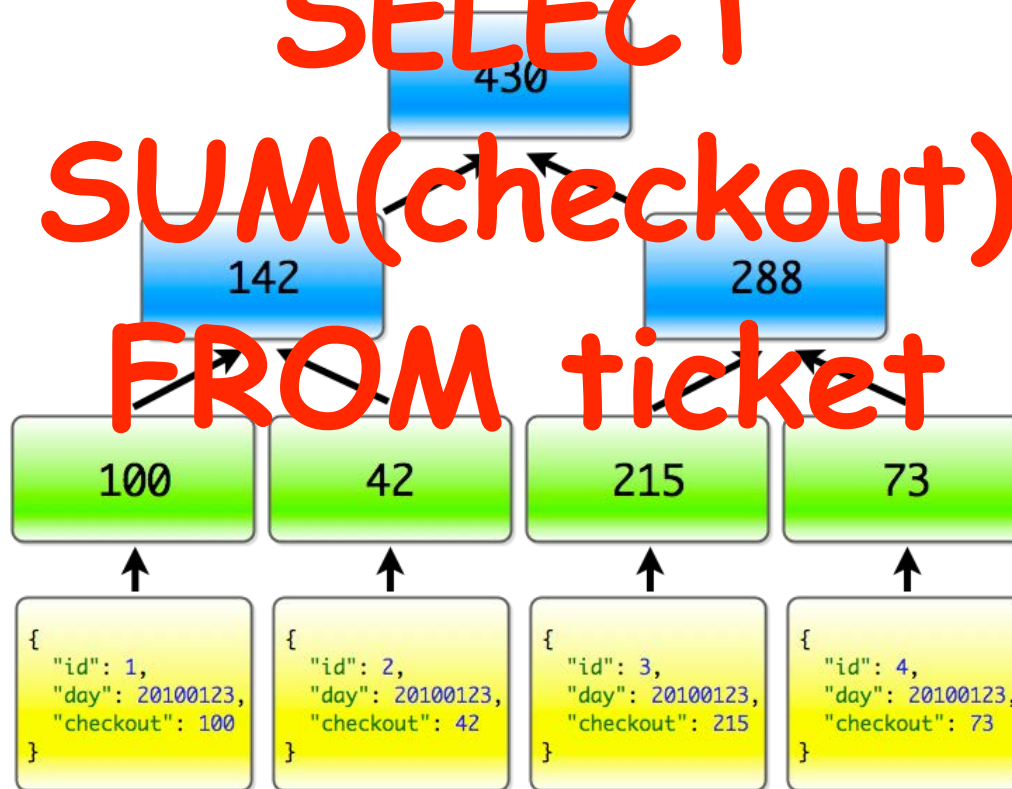
Reduce must be associative

$$\text{reduce}(\begin{array}{|c|c|c|c|} \hline 100 & 42 & 215 & 73 \\ \hline \end{array}) == \begin{array}{|c|} \hline 430 \\ \hline \end{array}$$


Must be equal to

$$\begin{array}{l} \text{reduce}(\text{reduce}(\begin{array}{|c|c|} \hline 100 & 42 \\ \hline \end{array}) == \begin{array}{|c|} \hline 142 \\ \hline \end{array}, \\ \text{reduce}(\begin{array}{|c|c|} \hline 215 & 73 \\ \hline \end{array}) == \begin{array}{|c|} \hline 288 \\ \hline \end{array}) \\) == \begin{array}{|c|} \hline 430 \\ \hline \end{array} \end{array}$$

**SELECT
SUM(checkout)
FROM ticket**



Inherently distributed



Calculate total checkout

#Aggregate alternative

```
db.tickets.aggregate ({
  "$group": {_id: null, "value": {$sum: "$checkout"}}},
  {$out: "sumOfCheckouts_agg"
})
```

#Map-reduce alternative

```
var map = function() {
  emit(null, this.checkout)
}
```

```
var reduce = function (key, values) {
  var sum=0
  for (var idx = 0; idx< values.length; idx++)
    sum+=values[idx];
  return sum;
}
```

```
db.tickets.mapReduce (map, reduce, {"out": "sumOfCheckouts"})
```

```
db.sumOfCheckouts.findOne().value
```

Persistent Collection

Sum(checkout) Group By day

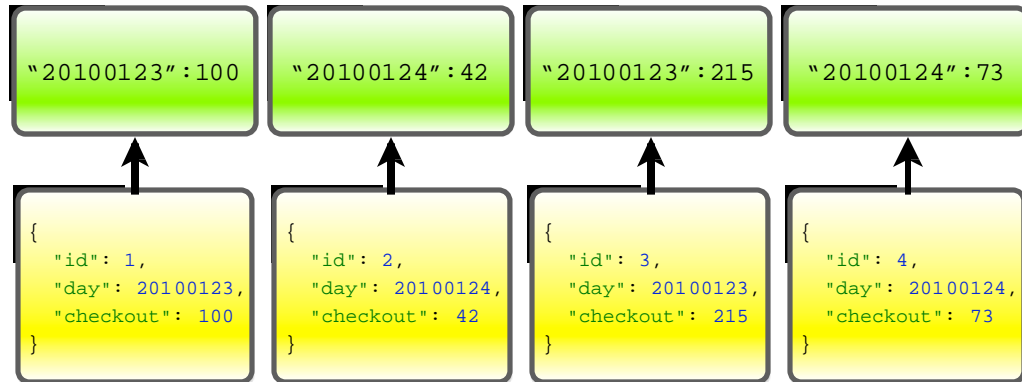
```
{  
  "id": 1,  
  "day": 20100123,  
  "checkout": 100  
}
```

```
{  
  "id": 2,  
  "day": 20100124,  
  "checkout": 42  
}
```

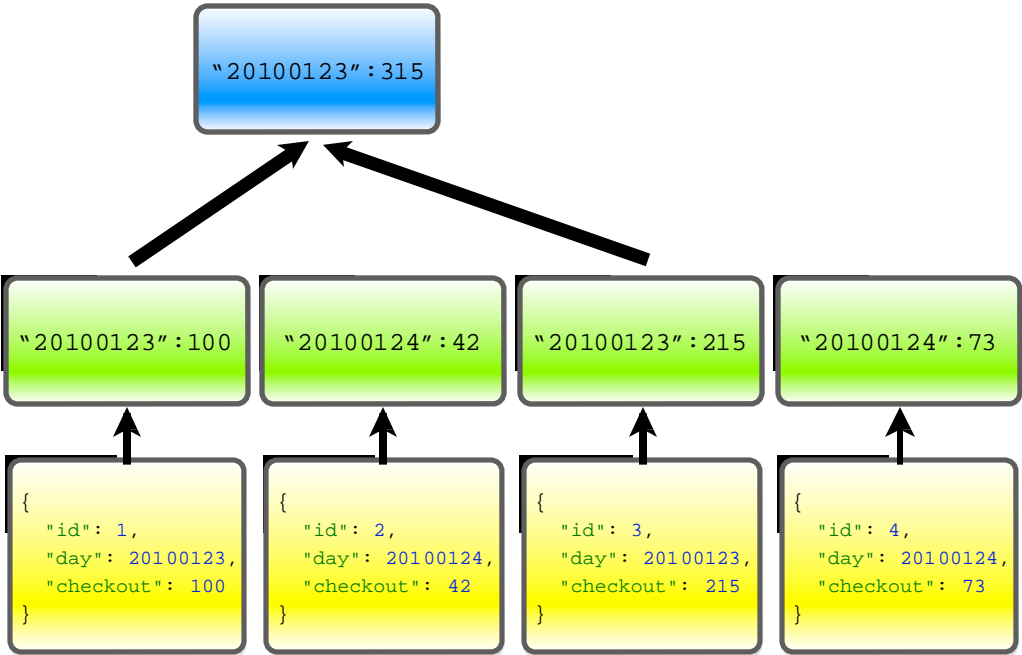
```
{  
  "id": 3,  
  "day": 20100123,  
  "checkout": 215  
}
```

```
{  
  "id": 4,  
  "day": 20100124,  
  "checkout": 73  
}
```

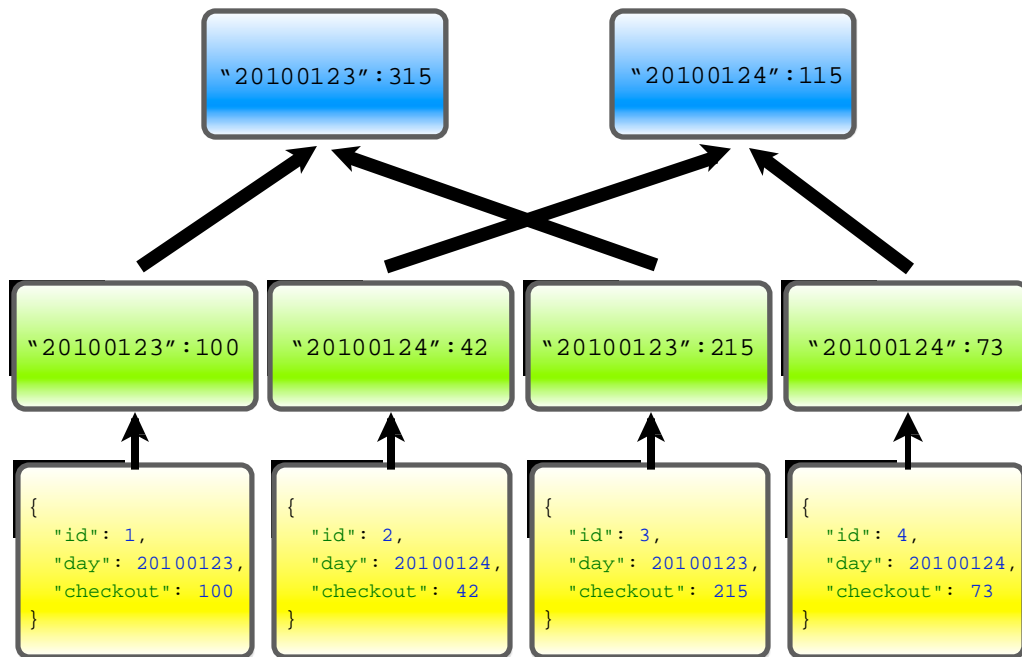
Map: emit(day,checkout)



Reduce: sum(checkouts)



Reduce: sum(checkouts)



Update the tickets to remove checkout

```
> db.tickets.update({ "_id": 1 }, {  
... $set: { "products": {  
..... "apple": { "qty":      5, "price": 10 },  
..... "kiwi": { "qty":      2, "price": 25 }  
..... }  
... },  
... $unset: { "checkout": 1 }  
... })  
  
> db.tickets.find()  
{ "_id" : 1, "day" : 20190123, "products" : {  
  "apple" : { "qty" : 5, "price" : 10 },  
  "kiwi" : { "qty" : 2, "price" : 25 }  
}}  
{ "_id" : 2, "day" : 20190123, "checkout" : 42 }  
{ "_id" : 3, "day" : 20190123, "checkout" : 215 }  
{ "_id" : 4, "day" : 20190123, "checkout" : 73 }
```


Sum(Checkout) by day Calculate Checkout

```
> var map = function() {  
... var checkout = 0  
... for (var name in this.products) {  
..... var product = this.products[name]  
..... checkout += product.qty * product.price  
..... }  
... emit(this.day, checkout)  
}  
  
> var reduce = function(key, values) {  
... var sum = 0  
... for (var index in values) sum += values[index]  
... return sum  
}
```

Sum(Checkout) by day Calculate Checkout

```
> db.tickets.mapReduce(map, reduce, { "out": "sumOfCheckouts" })

> db.sumOfCheckouts.find( )
{ "_id" : 20190123, "value" : 315 }
{ "_id" : 20190124, "value" : 110 }
```

For today

Follow the map-reduce examples linked below and create the orders collection:

<https://docs.mongodb.com/manual/tutorial/map-reduce-examples/>

Write a map-reduce aggregation that returns the number of items bought (the number of elements in the items array) per customer and submit it to ICON.