


# **Spatial and text indexes**

# Indexes

- MongoDB uses B-Tree indexes
- Can build the index on any field of the document
- Skips documents that do not have the indexed field (Sparse index)



Collection

# Examples

```
{ "_id": ObjectId(...),  
  "name": "John Doe",  
  "address": {  
    "street": "Main",  
    "zipcode": "53511",  
    "state": "WI"  
  }  
}
```

db.people.createIndex("name": 1)



**Field Level**

db.people.createIndex("address.zipcode": 1)



**Sub-Field Level**

db.people.createIndex("address": 1)



**Embedded document Level**  
(equality search only)

# Examples


```
{ "_id": ObjectId(...),  
  "name": "John Doe",  
  "address": {  
    "street": "Main",  
    "zipcode": "53511",  
    "state": "WI"  
  }  
}
```



**Compound-Field Index**

```
db.people.createIndex({"name": 1, "_id": -1})
```

```
db.people.find("_id": 1000)
```



**Index cannot answer this query  
(must have a predicate on "name")**


# Index Creation Options

```
{ "_id": ObjectId(...),  
  "name": "John Doe",  
  "address": {  
    "street": "Main",  
    "zipcode": "53511",  
    "state": "WI"  
  }  
}
```

```
db.people.createIndex({"name": 1, "_id": -1},  
                      {"background: True", "sparse": True,  
                       "unique": True})
```


# Geospatial indexes

[Longitude, Latitude]




# Surface type

## Flat



2d Indexes

## Spherical



2dsphere Indexes

# Quad Trees

- Split on *all* (two) dimensions at each level
- Split key space into equal size partitions (quadrants)
- Add a new node by adding to a leaf, and, if the leaf is already occupied, split until only one node per leaf

quadrant



quad tree node


Center:		Keys		Value	
		x		y	
Quadrants:		0,0	1,0	1,1	0,1

Center



# Quadtree


Simplest spatial structure on Earth !



# Quadtree


Simplest spatial structure on Earth !

321



# Quadtree


Simplest spatial structure on Earth !



# Quadtree

Simplest spatial structure on Earth !


32





# Quadtree

Simplest spatial structure on Earth !

321




# Quadtree Example



# Geospatial operators

**\$geoWithin**




No index, 2d, 2dsphere

**\$geoIntersects**



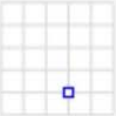
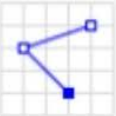


Index required - 2dsphere

**\$near/\$nearSphere**



2d, 2dsphere

# GeoJSON

Type	Examples	
Point		<pre>{ "type": "Point",   "coordinates": [30, 10] }</pre>
LineString		<pre>{ "type": "LineString",   "coordinates": [     [30, 10], [10, 30], [40, 40]   ] }</pre>
Polygon		<pre>{ "type": "Polygon",   "coordinates": [     [[30, 10], [40, 40], [20, 40], [10, 20], [30, 10]]   ] }</pre>
		<pre>{ "type": "Polygon",   "coordinates": [     [[35, 10], [45, 45], [15, 40], [10, 20], [35, 10]],     [[20, 30], [35, 35], [30, 20], [20, 30]]   ] }</pre>



# GeoJSON

Type	Examples	
MultiPoint		<pre>{   "type": "MultiPoint",   "coordinates": [     [10, 40], [40, 30], [20, 20], [30, 10]   ] }</pre>
MultiLineString		<pre>{   "type": "MultiLineString",   "coordinates": [     [[10, 10], [20, 20], [10, 40]],     [[40, 40], [30, 30], [40, 20], [30, 10]]   ] }</pre>
MultiPolygon		<pre>{   "type": "MultiPolygon",   "coordinates": [     [       [[30, 20], [45, 40], [10, 40], [30, 20]]     ],     [       [[15, 5], [40, 10], [10, 20], [5, 10], [15, 5]]     ]   ] }</pre>
		<pre>{   "type": "MultiPolygon",   "coordinates": [     [       [[40, 40], [20, 45], [45, 30], [40, 40]]     ],     [       [[20, 35], [10, 30], [10, 10], [30, 5], [45, 20], [20, 35]],       [[30, 20], [20, 15], [20, 25], [30, 20]]     ]   ] }</pre>

# 2d Query

```
db.<collection>.find( { <location field> :  
    { $geoWithin :  
        { $box|$polygon|$center : <coordinates>  
    } } } )
```

```
db.<collection>.find( { <location field> :  
    { $near : [ <x> , <y> ]  
    } } )
```

```
db.<collection>.find( { loc: [ <x> , <y> ] } )
```

# Geospatial indexing zips collection

- `db.zips.createIndex( {loc: "2d"} )`
- `db.zips.find ( {loc: { $geoWithin: { $box: [ [-73,42.5], [-72, 43] ] } } } )`
- `db.zips.find ( {loc: { $geoWithin: { $center: [ [-73,42.5], 10 ] } } } )`
- `db.zips.find ( {loc: { $near: [-73,42.5] } } )`

# Text Indexes

- Over fields that are strings or array of strings
- Index is used when using **\$text** search operator
- Only one index on the collection
  - But it can include multiple fields

```
db.collection.createIndex({content:"text"});
```

**One field**

```
db.collection.createIndex({subject:"text",content:"text"});
```

**Two fields**

```
db.collection.createIndex({"$**":"text"});
```

**All text fields**

# \$Text

Text search in mongoDB (Exact match)

Uses a text index and searches the indexed fields

```
{ $text: { $search: <string>, $language: <string> } }
```

```
db.articles.find( { $text: { $search: "coffee" } } )
```

**Search for “coffee” in the indexed field(s)**

```
db.articles.find( { $text: { $search: "bake coffee cake" } } )
```

**Apply “OR” semantics**

# \$Text

Text search in mongoDB

Uses a text index and searches the indexed fields

```
{ $text: { $search: <string>, $language: <string> } }
```

```
db.articles.find( { $text: { $search: "\"coffee cake\"" } } )
```

**Treated as one  
sentence**

```
db.articles.find( { $text: { $search: "bake coffee -cake" } } )
```

**“bake” or “coffee”  
but not “cake”**

# \$Text Score

\$Text returns a score for each matching document  
Score can be used in your query

```
db.articles.find(  
  { $text: { $search: "cake" } },  
  { score: { $meta: "textScore" } }  
).sort( { score: { $meta: "textScore" } } ).limit(3)
```

For regular expression match use **\$regex** operator

# City\_inspections examples

- `db.city_inspections.createIndex({"$**": "text"})`
- `db.city_inspections.find( { $text: { $search: "food deli" } } )`
- `db.city_inspections.find( { $text: { $search: "\"food deli\"" } } )`
- `db.city_inspections.find( { $text: { $search: "grocery -cigarette" } } )`
- `db.city_inspections.find(  
 { $text: { $search: "passed" } },  
 { score: { $meta: "textScore" } }  
)`.sort( { score: { \$meta: "textScore" } } ).limit(3)



# Collection Modeling

# Collection Modeling

Modeling multiple collections that reference each other

In Relational DBs → FK-PK Relationships

In MongoDB, two options:

**Referencing** between two collections

Use Id of one and put in the other

Very similar to FK-PK in Relational DBs

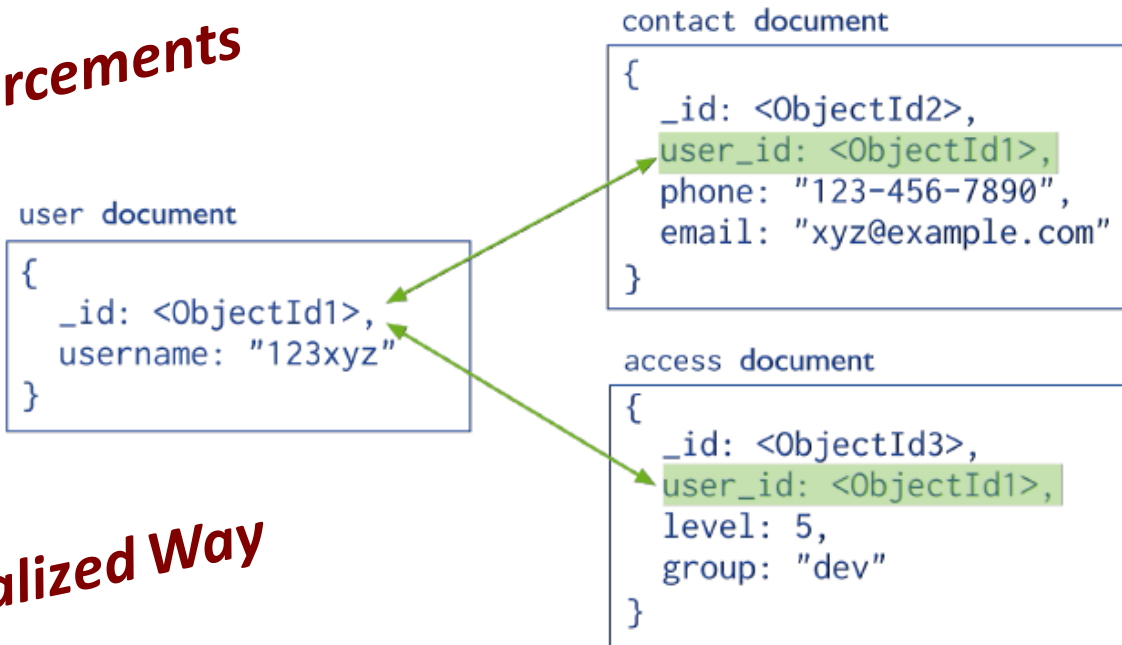
**Does not come with enforcement mechanism**

**Embedding** between two collections

Put the document from one collection inside the other one

# Referencing

**No Enforcements**



**Normalized Way**

- Have three collections in the DB: “User”, “Contact”, “Access”
- Link them by `_id` (or any other field(s))

# Embedding



**De-Normalized Way**

Have one collection in DB: "User"

The others are embedded inside each user's document

# Examples (1)

“Patron” & “Addresses”

```
{  
  _id: "joe",  
  name: "Joe Bookreader"  
}
```

```
{  
  patron_id: "joe",  
  street: "123 Fake Street",  
  city: "Faketon",  
  state: "MA",  
  zip: "12345"  
}
```

**Referencing**

- If it is 1-1 relationship
- If usually read the address with the name
- If address document usually does not expand

**If most of these hold  
→ better use Embedding**

# Examples (2)

“Patron” & “Addresses”

```
{
  _id: "joe",
  name: "Joe Bookreader",
  address: {
    street: "123 Fake Street",
    city: "Faketon",
    state: "MA",
    zip: "12345"
  }
}
```

**Embedding**

- When you read, you get the entire document at once
- In Referencing → Need to issue multiple queries

# Examples (3)

What if a “Patron” can have many “Addresses”

```
{
  _id: "joe",
  name: "Joe Bookreader"
}
```

```
{
  patron_id: "joe",
  street: "123 Fake Street",
  city: "Faketon",
  state: "MA",
  zip: "12345"
}
```

**Referencing**

- Do you read them together → Go for Embedding
- Are addresses dynamic (e.g., add new ones frequently)  
→ Go for Referencing

# Examples (4)

What if a “Patron” can have many “Addresses”

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

**Embedding**

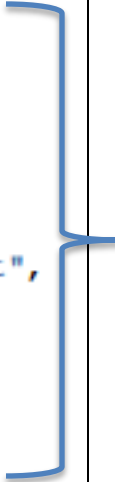
**Use array of addresses**



# Examples (5)

If addresses are added frequently ...

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```



This array will expand frequently



Size of "Patron" document increases frequently



May trigger re-locating the document each time (*Bad*)

# Document Size and Storage

Each document needs to be contiguous on disk

If doc size increases → Document location must change

If doc location changes → Indexes must be updated  
→ leads to more expensive updates

```
{
  _id: "joe",
  name: "Joe Bookreader",
  addresses: [
    {
      street: "123 Fake Street",
      city: "Faketon",
      state: "MA",
      zip: "12345"
    },
    {
      street: "1 Some Other Street",
      city: "Boston",
      state: "MA",
      zip: "12345"
    }
  ]
}
```

- Each document is allocated a **power-of-2 bytes** (the smallest above its size)
- Meaning, the system keeps some space empty for possible expansion

# Examples (6)

## One-to-Many “Book”, “Publisher”

A book has one publisher

A publisher publishes many books

### If embed “Publisher” inside “Book”

Repeating publisher info inside each of its books

Very hard to update publisher’s info

### If embed “Book” inside “Publisher”

Book becomes an array (many)

Frequently update and increases in size

*Referencing is better  
in this case*

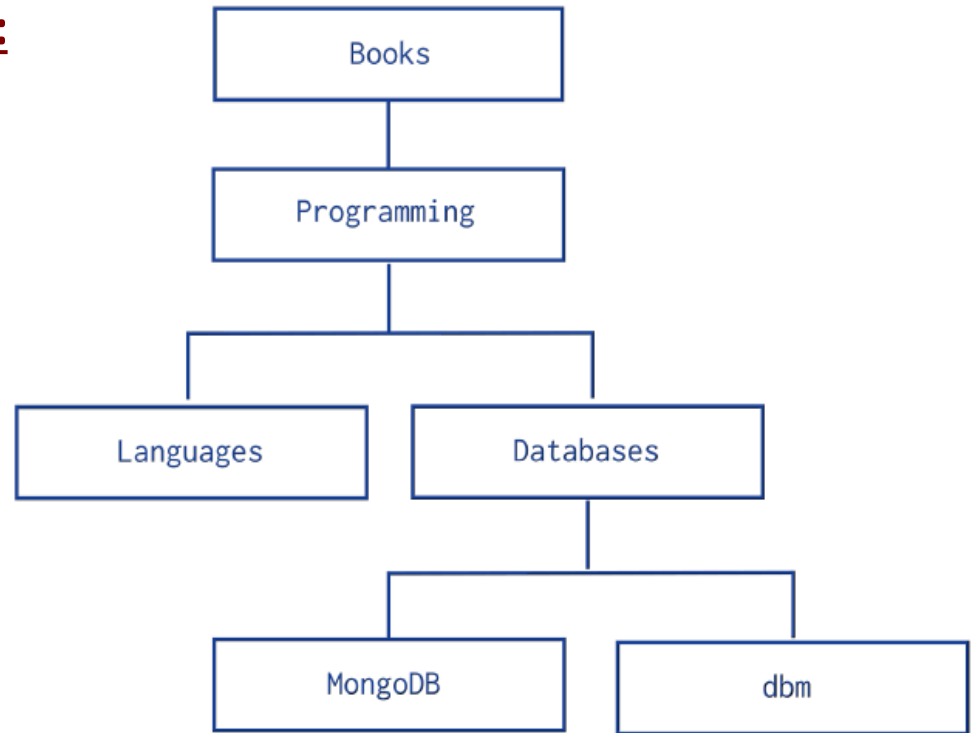
# Modeling Tree Structure

# Collections with Tree-Like Relationships

- Insert these records while maintaining this tree-like relationship

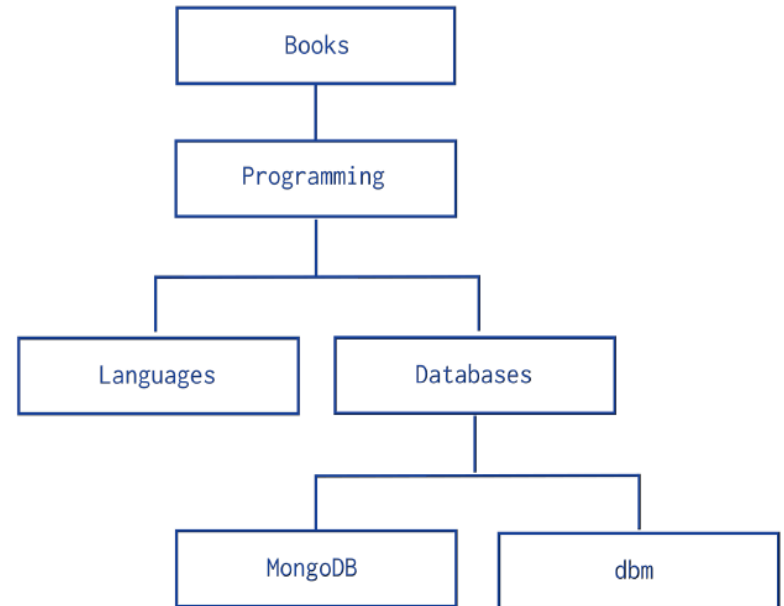
## Given one node, answer queries:

- Report the parent node
- Report the children nodes
- Report the ancestors
- Report the descendants
- Report the siblings



# Method 1: Parent References

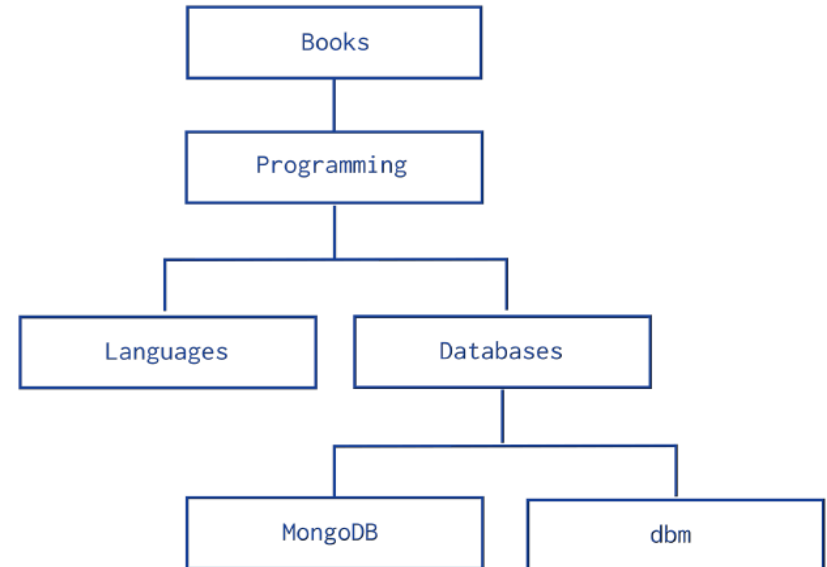
Each document has a field “parent”  
Order does not matter



```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )  
db.categories.insert( { _id: "dbm", parent: "Databases" } )  
db.categories.insert( { _id: "Databases", parent: "Programming" } )  
db.categories.insert( { _id: "Languages", parent: "Programming" } )  
db.categories.insert( { _id: "Programming", parent: "Books" } )  
db.categories.insert( { _id: "Books", parent: null } )
```

# Method 2: Child References

Each document has an array of immediate children




```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

# Method 1: Parent References

Q1: Parent of “Programming”

Q2: Siblings of “Databases”



```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )
db.categories.insert( { _id: "dbm", parent: "Databases" } )
db.categories.insert( { _id: "Databases", parent: "Programming" } )
db.categories.insert( { _id: "Languages", parent: "Programming" } )
db.categories.insert( { _id: "Programming", parent: "Books" } )
db.categories.insert( { _id: "Books", parent: null } )
```



# Method 1: Parent References

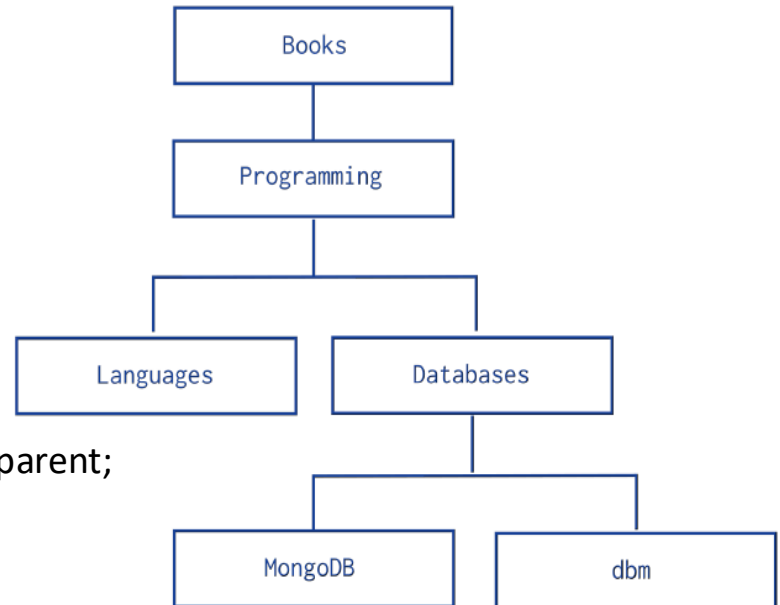
## Q1: Parent of “Programming”

```
db.categories.find( { _id: "Programming" }, { parent: 1, _id: 0 } );
```

## Q2: Siblings of “Databases”

```
var parentDoc = db.categories.findOne( { _id: "Databases" } ).parent;
```

```
db.categories.find( { parent: parentDoc,
                     _id: { $ne: "Databases" } } );
```

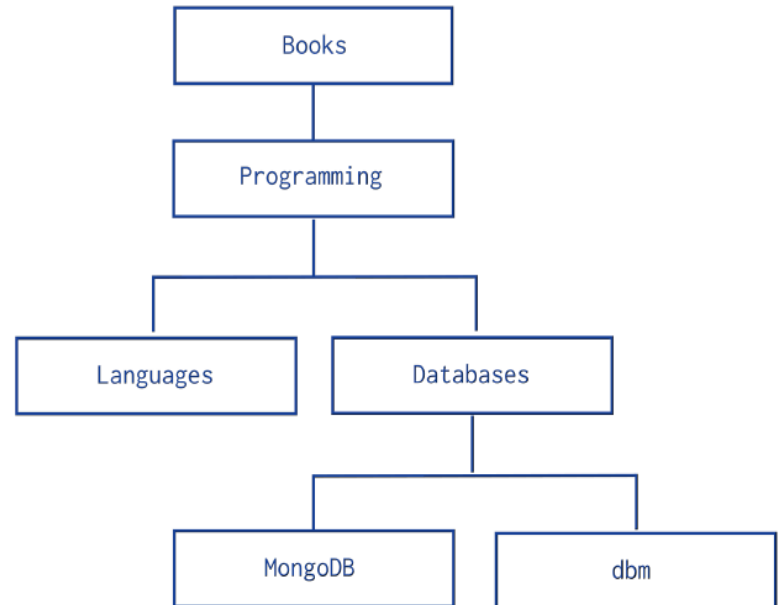


```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )
db.categories.insert( { _id: "dbm", parent: "Databases" } )
db.categories.insert( { _id: "Databases", parent: "Programming" } )
db.categories.insert( { _id: "Languages", parent: "Programming" } )
db.categories.insert( { _id: "Programming", parent: "Books" } )
db.categories.insert( { _id: "Books", parent: null } )
```

# Method 1: Parent References

Q3: Descendants of “Programming”

Complex...Requires  
recursive calls




```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )
db.categories.insert( { _id: "dbm", parent: "Databases" } )
db.categories.insert( { _id: "Databases", parent: "Programming" } )
db.categories.insert( { _id: "Languages", parent: "Programming" } )
db.categories.insert( { _id: "Programming", parent: "Books" } )
db.categories.insert( { _id: "Books", parent: null } )
```

# Method 1: Parent References

## Q3: Descendants of “Programming”

```
var descendants = [];  
var stack = [];  
var item = db.categories.findOne({_id: "Programming"});  
stack.push(item);  
while (stack.length > 0) {  
    var current = stack.pop();  
    var children = db.categories.find({parent: current._id});  
    while (children.hasNext() == true) {  
        var child = children.next();  
        descendants.push(child._id);  
        stack.push(child);  
    }  
}  
descendants;
```




# Method 1: Parent References

## Q4: Ancestors of “MongoDB”

Try it yourself....

Should be:


“Databases”, “Programming”, “Books”



```
db.categories.insert( { _id: "MongoDB", parent: "Databases" } )
db.categories.insert( { _id: "dbm", parent: "Databases" } )
db.categories.insert( { _id: "Databases", parent: "Programming" } )
db.categories.insert( { _id: "Languages", parent: "Programming" } )
db.categories.insert( { _id: "Programming", parent: "Books" } )
db.categories.insert( { _id: "Books", parent: null } )
```

# Method 2: Child References

**Q1: Get children documents of “Programming”**




```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

# Method 2: Child References

## Q1: Get children documents of “Programming”

```
var x = db.categories.findOne({_id: "Programming"}).children;
```

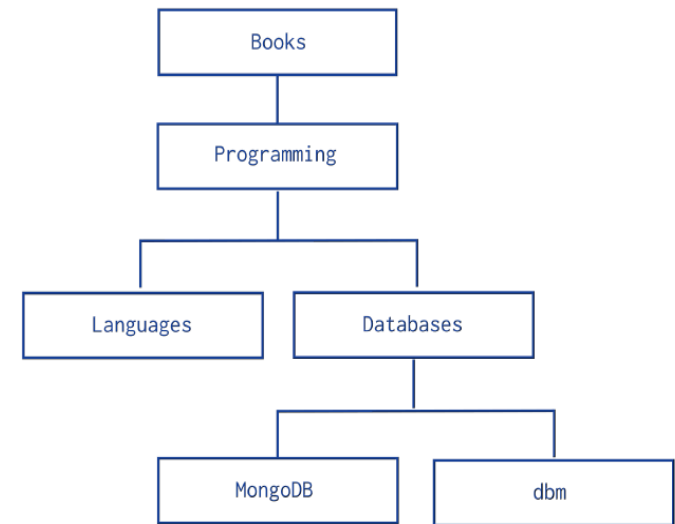
```
db.categories.find({_id: {$in: x}});
```



```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

# Method 2: Child References

## Q2: Ancestors of “MongoDB”



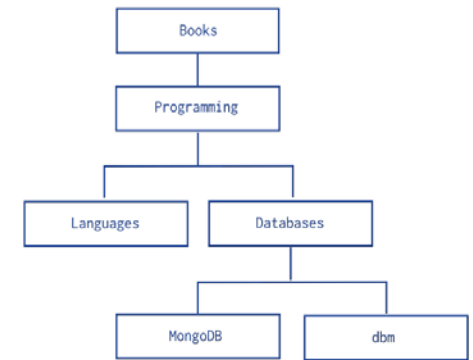
```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

# Method 2: Child References

## Q2: Ancestors of “MongoDB”

```
var results=[];
var parent = db.categories.findOne({children: "MongoDB"});
while(parent){
    print({Message: "Going up one level..."});
    results.push(parent._id);
    parent = db.categories.findOne({children: parent._id});
}

results;
```






# Method 2: Child References

Q3: descendants of “Books”

Try it yourself....

Should be all nodes



```
db.categories.insert( { _id: "MongoDB", children: [] } )
db.categories.insert( { _id: "dbm", children: [] } )
db.categories.insert( { _id: "Databases", children: [ "MongoDB", "dbm" ] } )
db.categories.insert( { _id: "Languages", children: [] } )
db.categories.insert( { _id: "Programming", children: [ "Databases", "Languages" ] } )
db.categories.insert( { _id: "Books", children: [ "Programming" ] } )
```

# For today

- Download and install Neo4j community edition
- <https://neo4j.com/download-center/#community>

If you don't already have it, you will also need to install Oracle JDK or Open JDK - Java Development Kit Standard Edition.

Visit <http://localhost:7474> in your web browser  
Default username and password: 'neo4j'

Submit a screenshot to ICON