

Graph Databases

Data is more connected:

- Text
- HyperText
- RSS
- Blogs
- Tagging
- RDF

Data is more Semi-Structured:

- If you tried to collect all the data of every movie ever made, how would you model it?
- Actors, Characters, Locations, Dates, Costs, Ratings, Showings, Ticket Sales, etc.



Document Databases

- MongoDB, CouchDB
- Pros:
 - Simple, powerful data model
 - Scalable
- Cons
 - Poor for interconnected data
 - Query model limited to keys and indexes
 - Map reduce for larger queries

Graph Databases

- Data Model:
 - Nodes and Relationships
- Examples:
 - Neo4j, OrientDB, InfiniteGraph, AllegroGraph
- Cypher
 - Graph database query language

Graph Databases: Pros and Cons

- Pros:
 - Powerful data model, as general as RDBMS
 - Connected data locally indexed
 - Easy to query
- Cons
 - Sharding (lots of people working on this)
 - Scales UP reasonably well
 - Requires rewiring your brain

What are graphs good for?

- Real Time Recommendations
- Master Data Management
- Fraud Detection
- Social computing
- Systems management
- Web of things
- Genealogy
- Time series data
- Product catalogue
- Web analytics
- Scientific computing (especially bioinformatics)
- And much more!


What is a Graph?

What is a Graph?

- An abstract representation of a set of objects where some pairs are connected by links.




Object (Vertex, Node)




Link (Edge, Arc, Relationship)

Different Kinds of Graphs

- Undirected Graph




- Directed Graph




- Pseudo Graph



- Multi Graph




- Hyper Graph



More Kinds of Graphs


- Weighted Graph



- Labeled Graph




- Property Graph




What is a Graph Database?

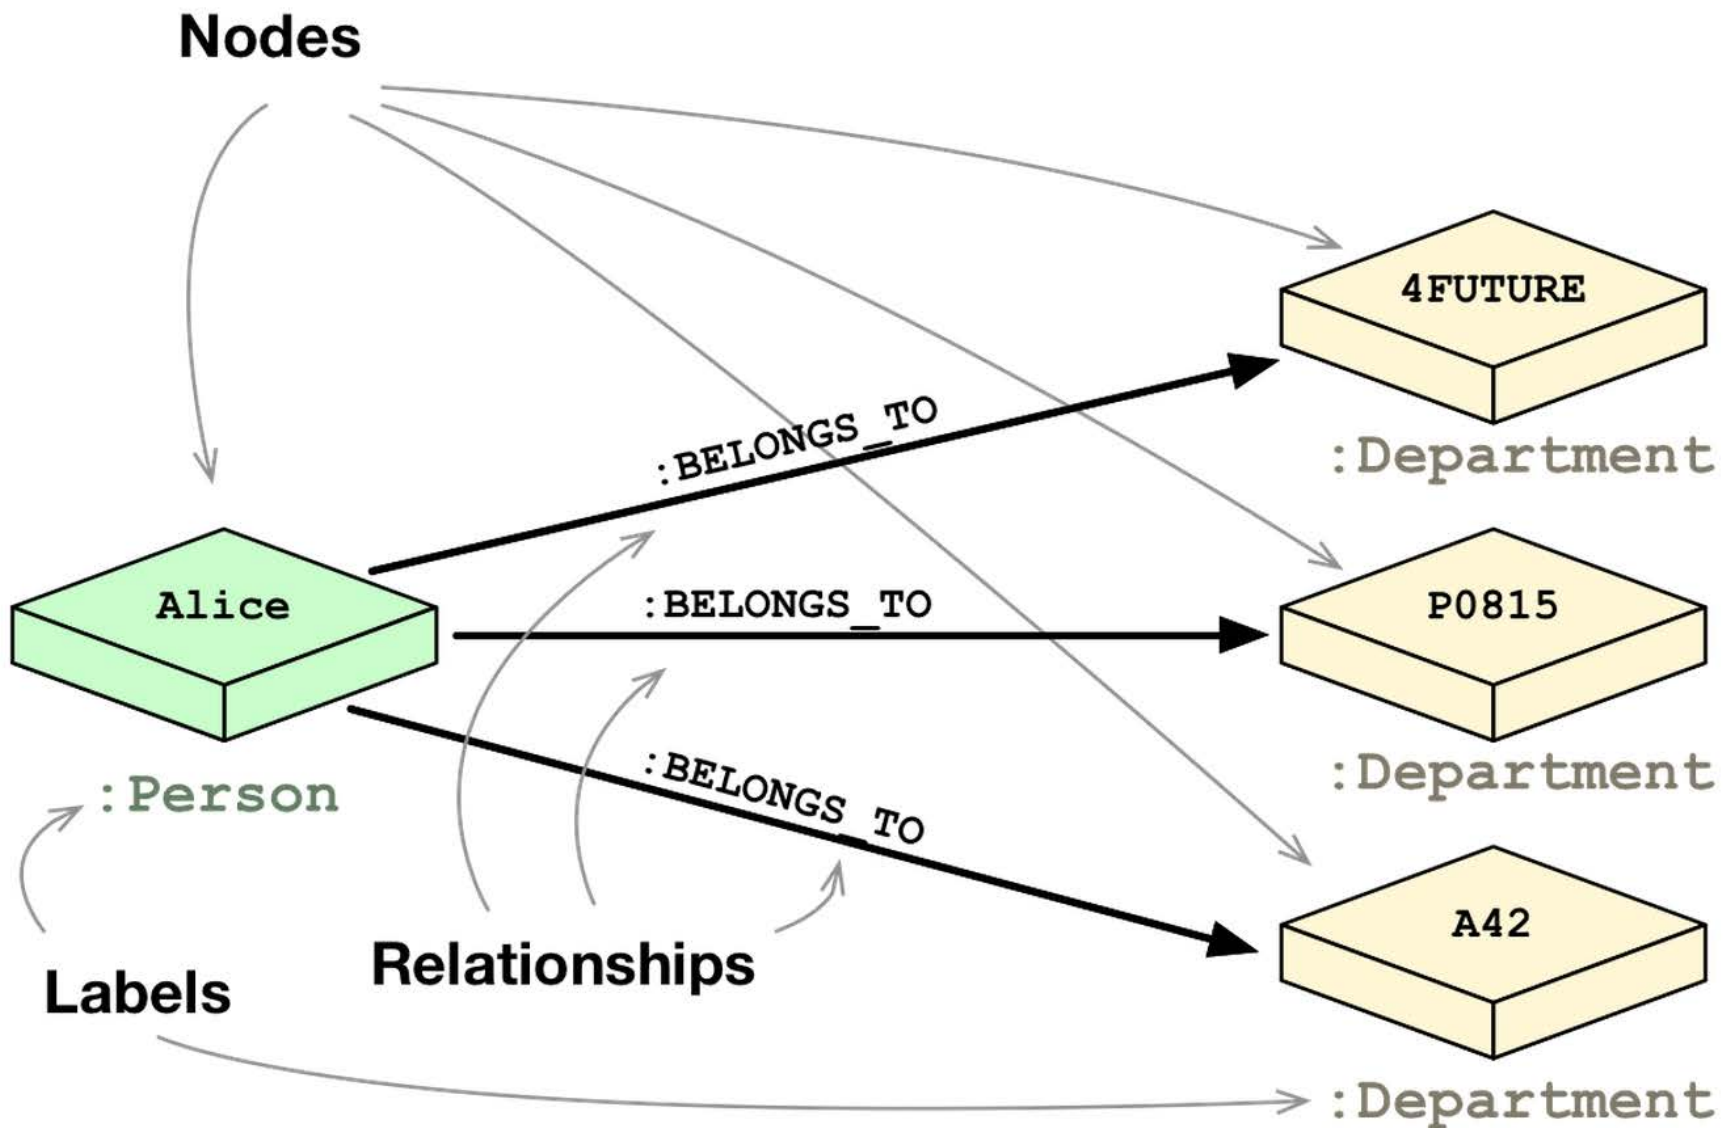
- A database with an explicit graph structure
- Each node knows its adjacent nodes
- As the number of nodes increases, the cost of a local step (or hop) remains the same
- Plus an Index for lookups

Relational Databases

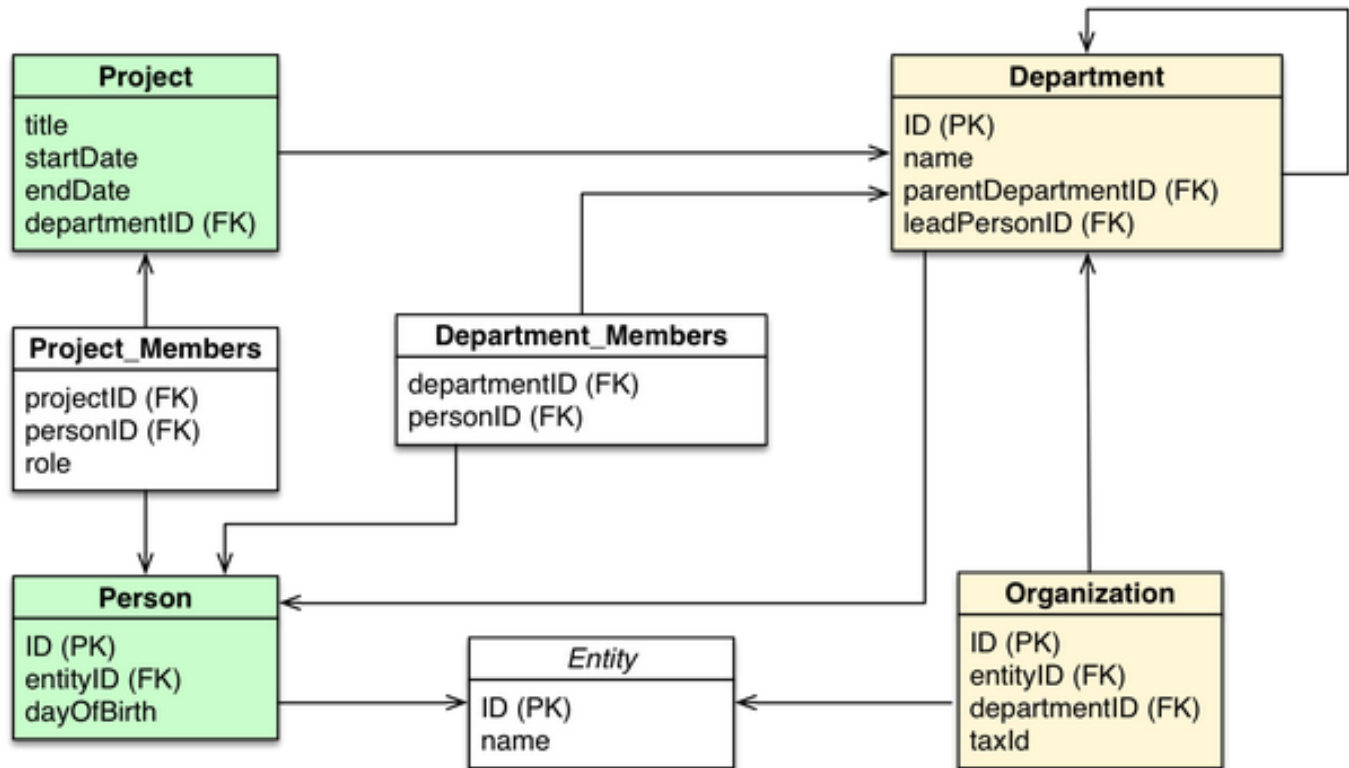


Graph Databases

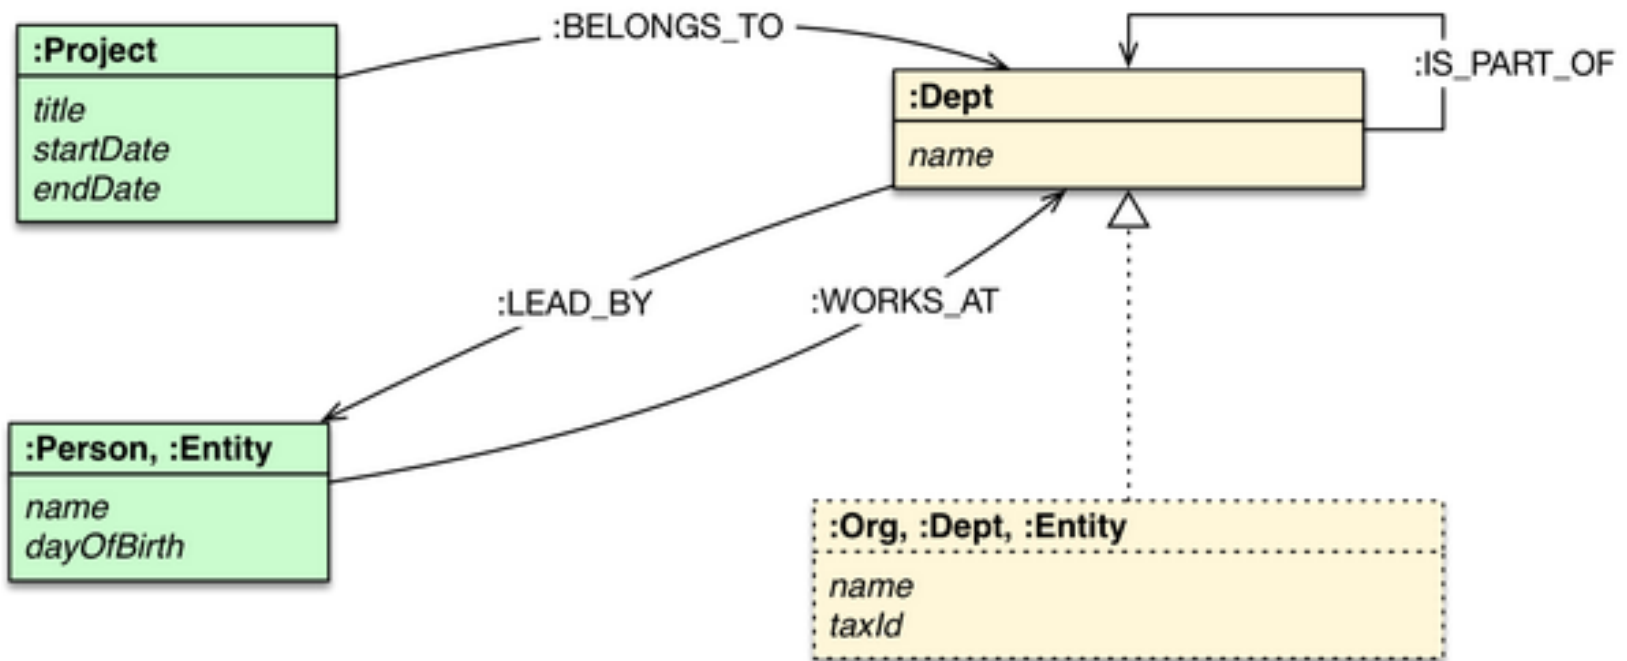




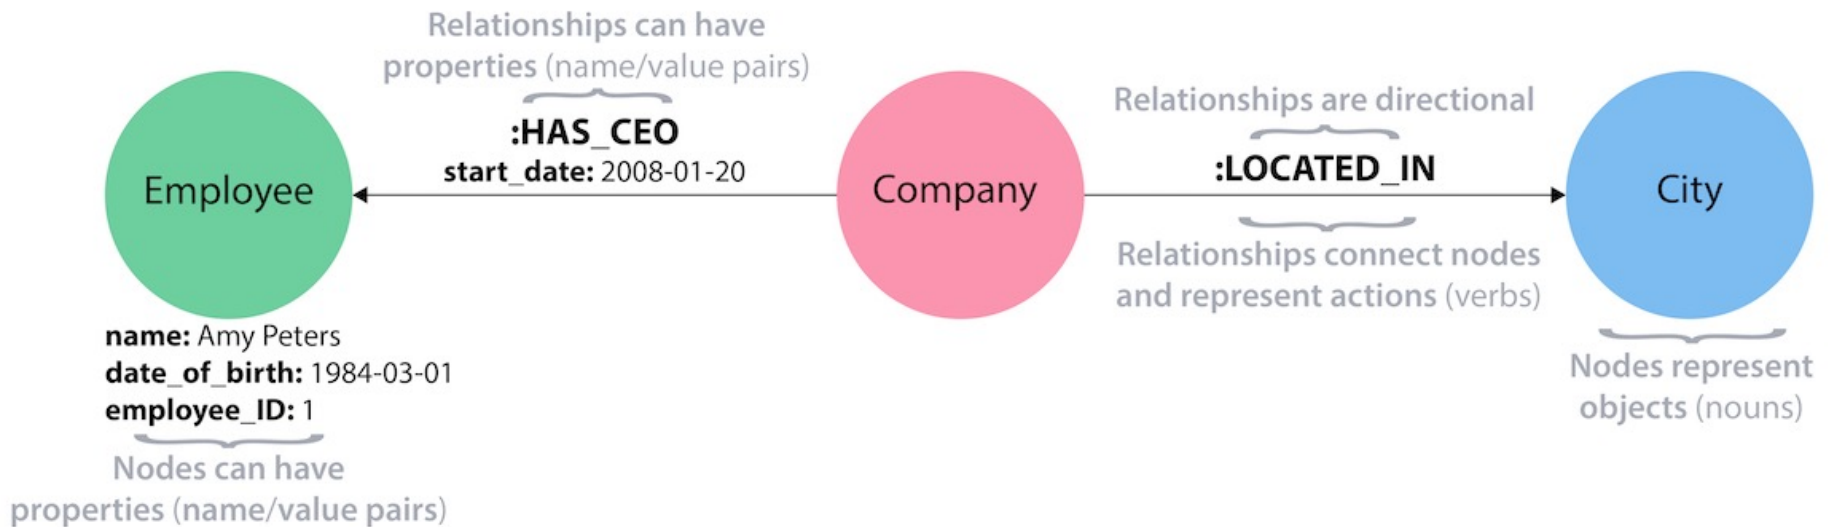
Relational Databases



Graph Databases



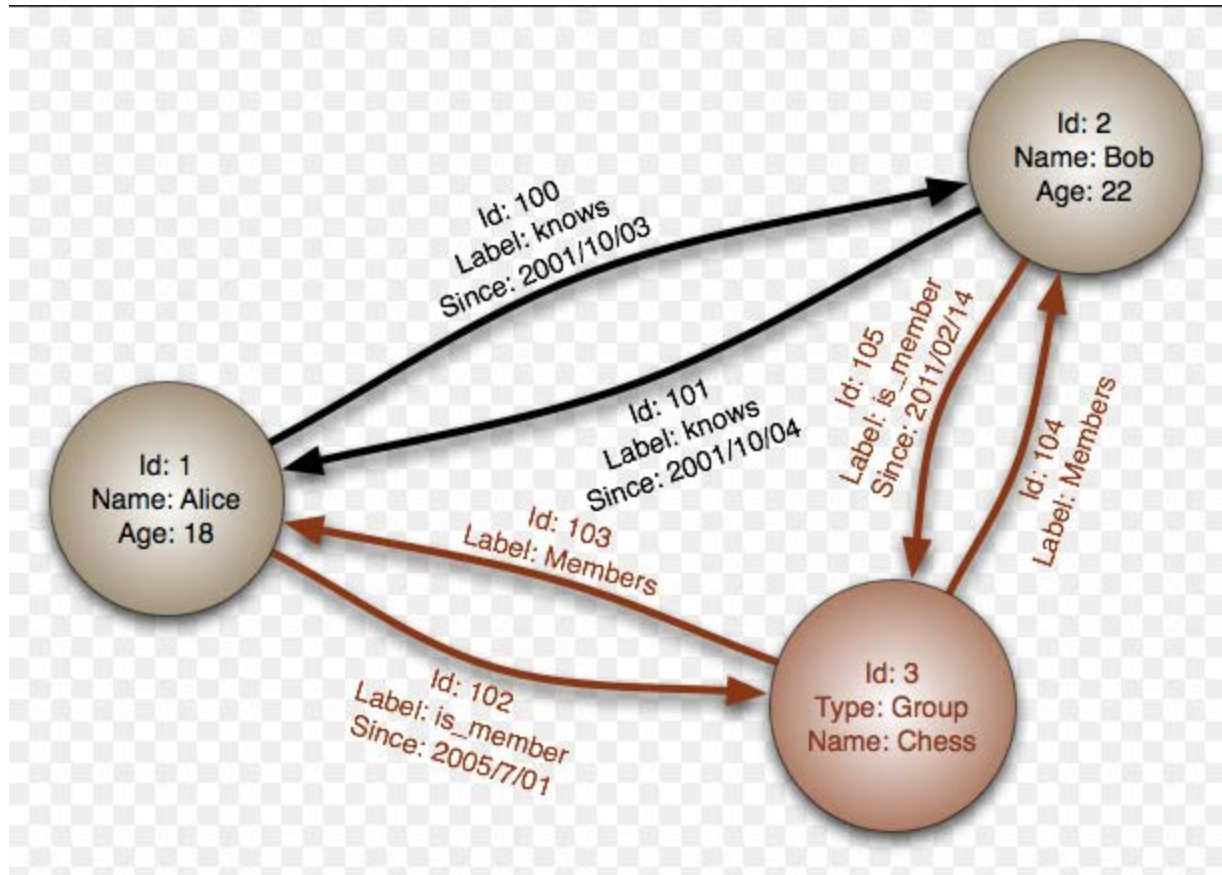
Property graph model



Graph Databases

- Database that uses graph structures with nodes, edges and properties to store data
- Provides index-free adjacency
 - Every node is a pointer to its adjacent element
- Edges hold most of the important information and connect
 - nodes to other nodes
 - nodes to properties

Graph Databases



Advantage of Graph Databases

- When there are relationships that you want to analyze Graph databases become a very nice fit because of the data structure
- Graph databases are very fast for associative data sets
 - Like social networks
- Map more directly to object oriented applications
 - Object classification and Parent->Child relationships

Disadvantages

- If data is just tabular with not much relationship between the data, graph databases do not fare well
- OLAP support for graph databases is not well developed
 - Lots of research happening in this area





Diagram illustrating a table structure with labels:

- Table name: STUDENTS
- Column name: Rollno, Name, Phone
- Tuple / Row: s1, s2, s3, s4
- Attribute / Column: Rollno, Name, Phone
- Table / Relation: The entire table structure


Rollno	Name	Phone
s1	Louis Figo	454333
s2	Raul	656675
s3	Roberto Carlos	546782
s4	Guti	567345



Ease of aggregation

What is Neo4j

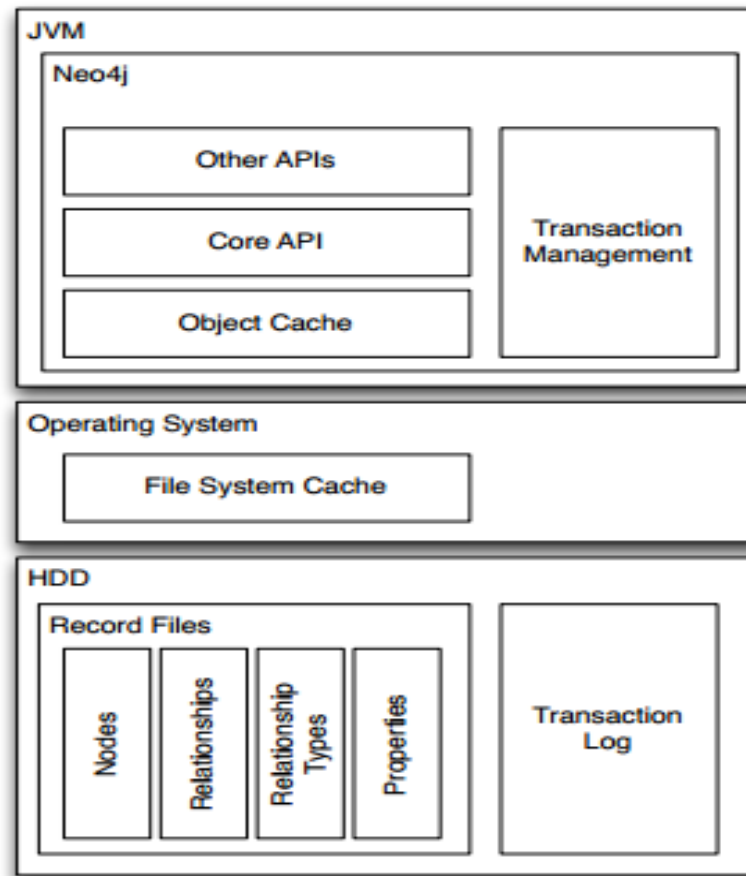
- Developed by Neo Technologies
- Most Popular Graph Database
- Implemented in Java
- Open Source



Salient features of Neo4j

- Neo4j is schema free – Data does not have to adhere to any convention
- ACID – atomic, consistent, isolated and durable for logical units of work
- Easy to get started and use
- Well documented and large developer community
- Support for wide variety of languages
 - Java, Python, Perl, Scala, Cypher, etc

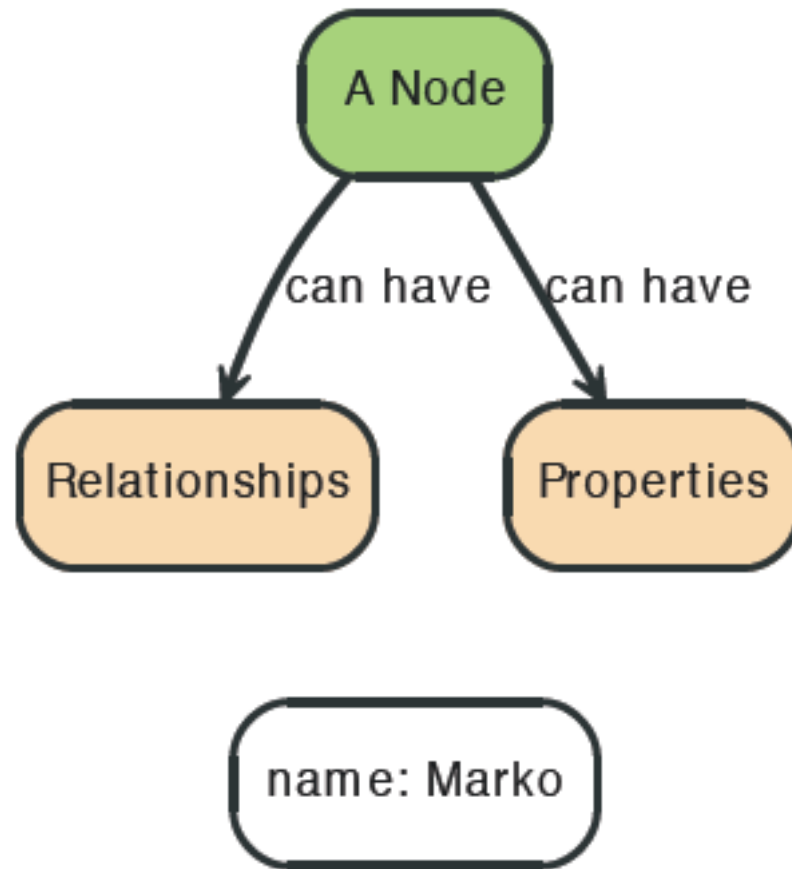
Neo4j Software Architecture



Neo4j Tips

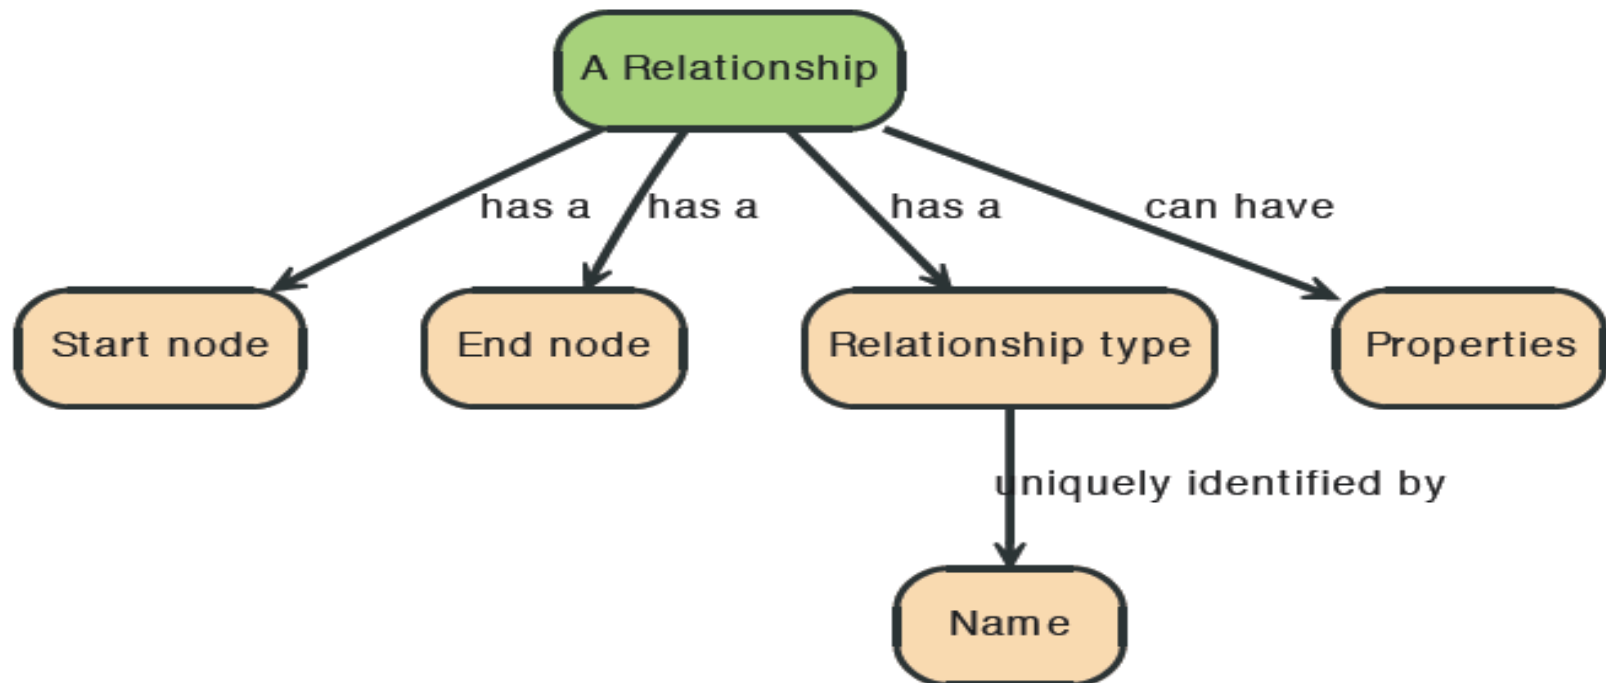
- Each entity table is represented by a label on nodes
- Each row in a entity table is a node
- Columns on those tables become node properties.
- Join tables are transformed into relationships, columns on those tables become relationship properties

Node in Neo4j





Relationships in Neo4j


- Relationships between nodes are a key part of Neo4j.



Relationships in Neo4j



Twitter and relationships

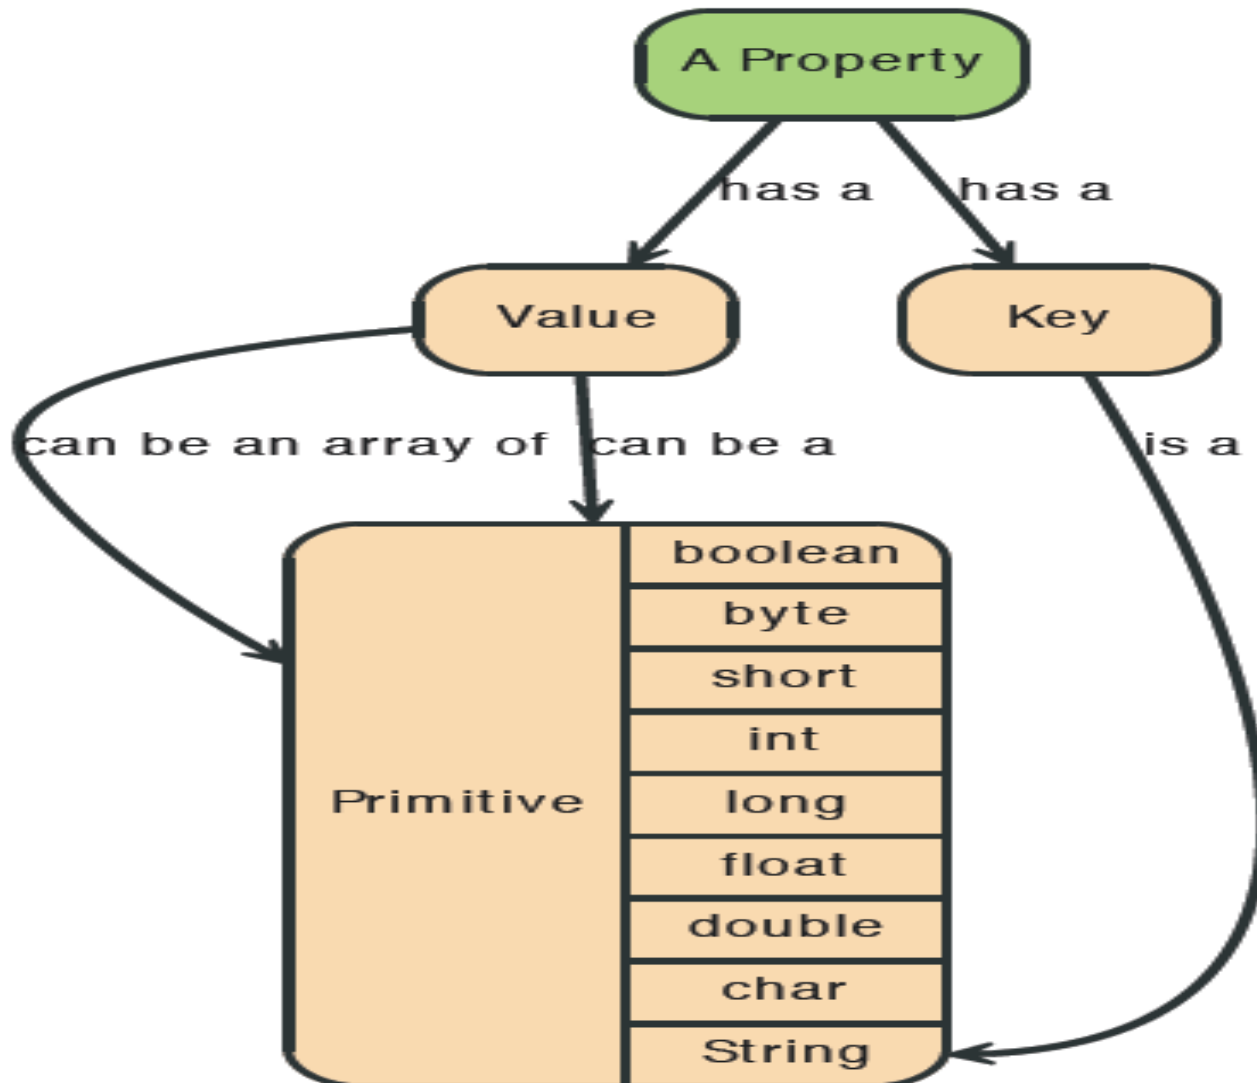


Properties

- Both nodes and relationships can have properties.
- Properties are key-value pairs where the key is a string.
- Property values can be either a primitive or an array of one primitive type.



For example `String`, `int` and `int[]` values are valid for properties.

Properties



Paths in Neo4j

- A path is one or more nodes with connecting relationships, typically retrieved as a query or traversal result.



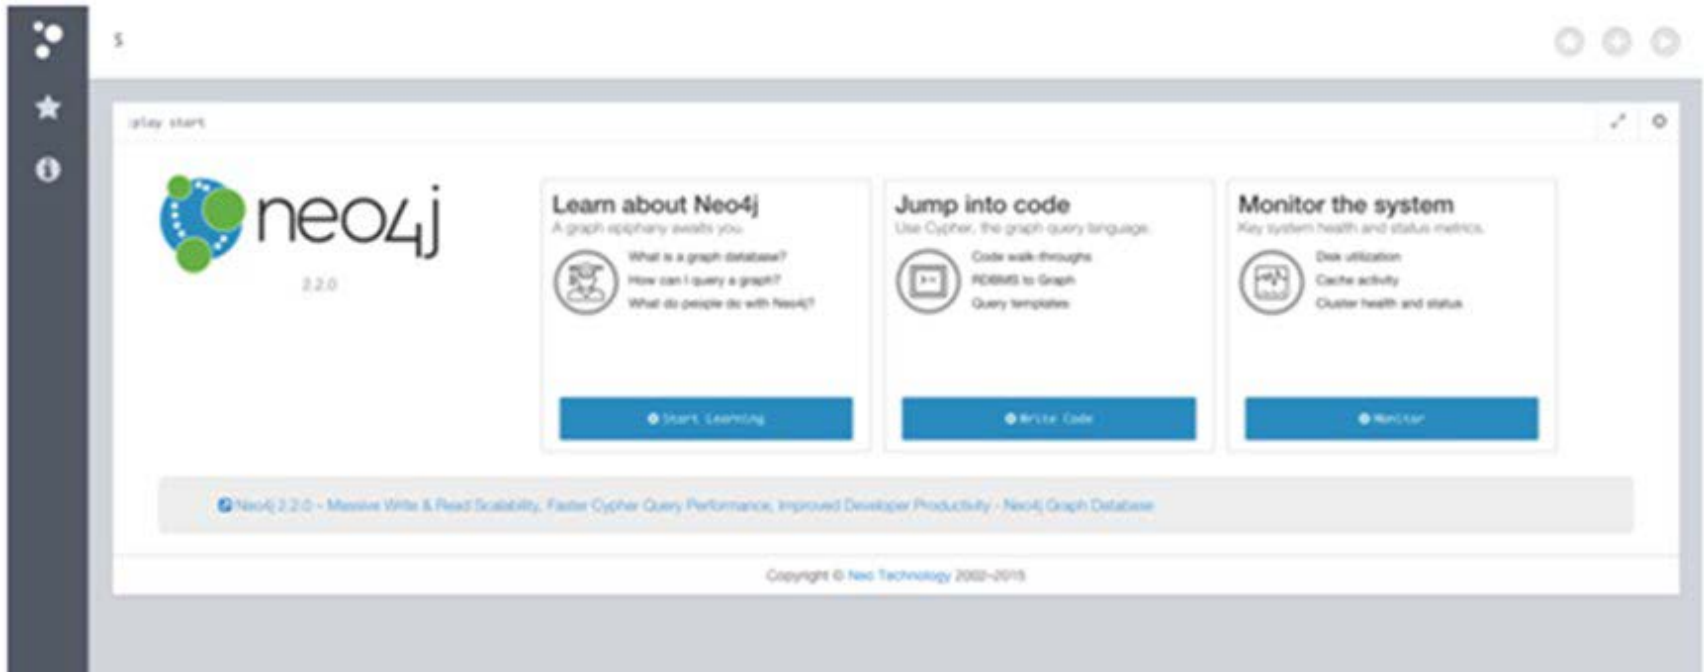
Neo4j browser

After you've downloaded and unzipped the Neo4j package, cd into the Neo4j directory and start up the server like this:

```
$ bin/neo4j start
```

To make sure you're up and running, try curling this URL:

```
$ curl http://localhost:7474/db/data/
```



Cypher

- Query Language for Neo4j
- Easy to formulate queries based on relationships
- Many features stem from improving on pain points with SQL such as join tables

Cypher basic syntax

- `()` Node
- `{}` Properties
- `[]` Relationships

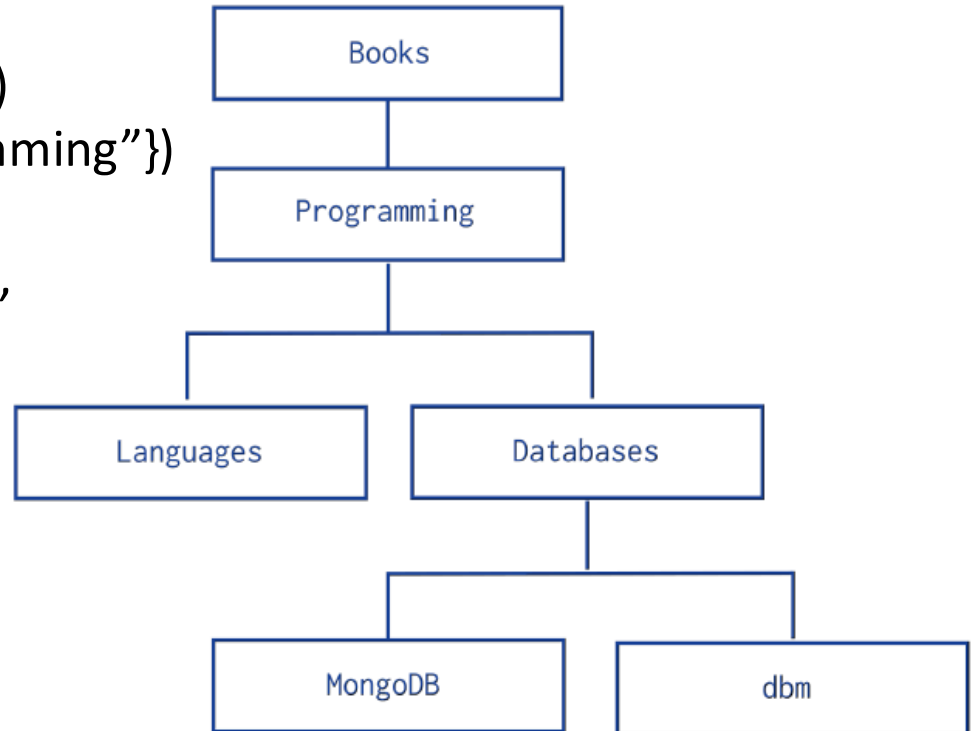
- `MATCH (n) RETURN n;`
- Returns every node in the database

- `MATCH (a)-[:CONNECTED]->(b) RETURN a, b;`
- Returns all nodes related with the `CONNECTED` relationship

Collections with Tree-Like Relationships

```
Create (c1:Category {name: "Books"})  
Create (c2:Category {name: "Programming"})
```


```
MATCH (c:Category {name: "Books"}),  
      (d:Category {name: "Programming"})  
CREATE (c)-[p:PARENT_OF]->(d)
```



Collections with Tree-Like Relationships

Given one node, answer queries:

- Report the parent node
Parent of “Programming”
- Report the children nodes
Children of “Programming”
- Report the ancestors
Ancestors of “MongoDB”
- Report the descendants
Descendants of “Programming”
- Report the siblings
Siblings of “Databases”



Categories example

Children of “Programming”

```
MATCH (p:Category {name: “Programming”}-[:PARENT_OF]->(c)
RETURN p,c
```

Parent of “Programming”

```
MATCH (p:Category {name: “Programming”}<-[:PARENT_OF]- (c)
RETURN p,c
```

Ancestors of “MongoDB”

```
MATCH (p:Category)-[:PARENT_OF*]->(c:Category{name: “MongoDB”})
RETURN p,c
```

Descendants of “Programming”

```
MATCH (p:Category {name: “Programming”}-[:PARENT_OF*]->(c)
RETURN p,c
```

Siblings of “Databases”

```
MATCH (p:Category {name: "Databases"})<-[:PARENT_OF]-(c),
(c:Category)-[:PARENT_OF]->(d)
RETURN d
```

For today...

- Create three nodes:
 - One for you and two other friends
 - Connect them with a FRIEND relationship
 - Query the three nodes, and submit a png of the graph view to ICON