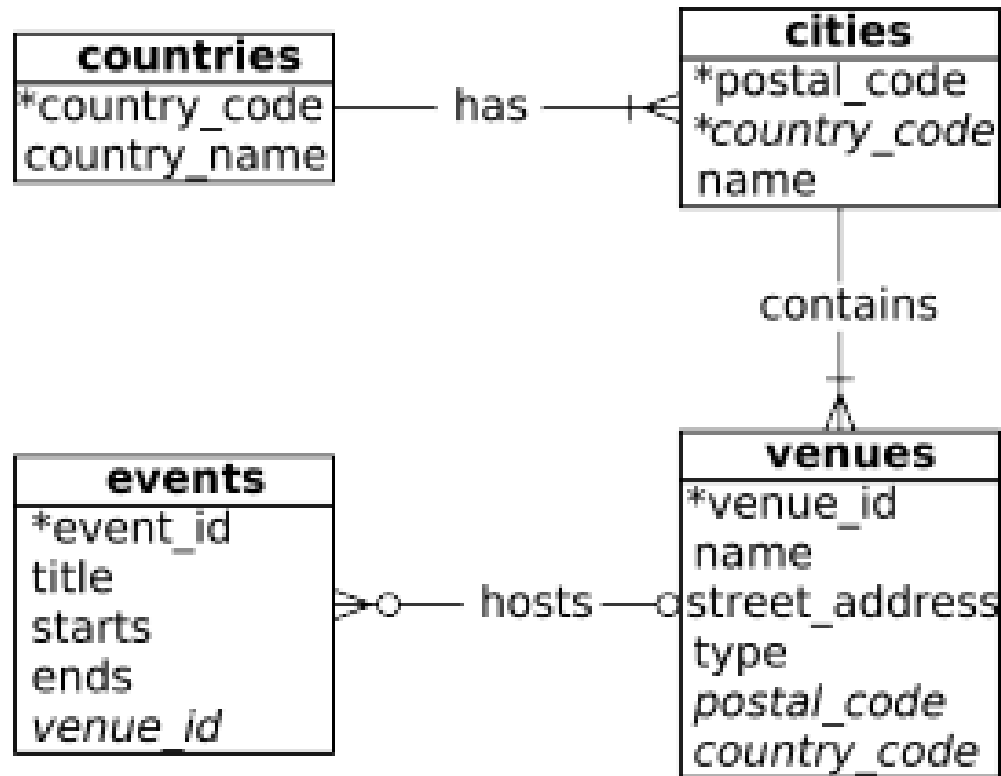


QUERIES and AGGREGATES

See <http://sqlzoo.net/>

Our working example



Let's add more venues and events.

Basic SQL Query



- Semantics/Conceptual evaluation strategy:
 - Compute the cross-product of *relation-list*.
 - Discard resulting tuples if they fail *qualifications*.
 - Delete attributes that are not in *target-list*.
 - If DISTINCT is specified, eliminate duplicate rows.
- Not an efficient evaluation plan! (Optimizer picks efficient plans)

Find venues with events on '19-Oct-2019'

```
SELECT v.name as Venue  
      FROM events e, venues v  
      WHERE e.venue_id = v.venue_id  
      AND (DATE(e.starts) = '2019-10-19')
```

- Add DISTINCT to this query. Effect?
- Equivalent SQL using JOIN?

Expressions and Strings

```
SELECT e.title,  
       to_char(e.starts, 'FMDay, Mon FMDD YYYY HH12:MI AM') as starts  
FROM events e  
WHERE e.title LIKE '%visit%'  
ORDER BY e.title;
```

- Illustrates date formatting and string pattern matching
- `AS` is a way to name fields in result
- `LIKE` is used for string matching. ``_`` stands for any one character and ``%`` stands for 0 or more arbitrary characters.

Find venues with events either in '22-Apr-2020' or '23-Apr-2020'

- UNION: Compute the union of two *union-compatible* sets of tuples
 - Same number/types of fields.
- Also available: INTERSECT and EXCEPT (What do we get if we replace UNION by EXCEPT?)
- SQL oddities: duplicates with union, except, intersect
 - Default: eliminate duplicates!
 - Use ALL to keep duplicates

```
SELECT v.name as Venue, v.country_code
FROM events e, venues v
WHERE e.venue_id = v.venue_id AND
(DATE(e.starts) = '2020-04-22' OR
DATE(e.starts) = '2020-04-23')
```

```
SELECT v.name as Venue, v.country_code
FROM events e, venues v
WHERE e.venue_id = v.venue_id
AND DATE(e.starts) = '2020-04-22'
```

UNION

```
SELECT v.name as Venue, v.country_code
FROM events e, venues v
WHERE e.venue_id = v.venue_id
AND DATE(e.starts) = '2020-04-23'
```

Find venues with events either in '22-Apr-2020' AND '23-Apr-2020'

- INTERSECT: Compute the intersection of any two *union-compatible* sets of tuples.
- In the SQL/92 standard, but some systems don't support it.

```
SELECT v.name as Venue, v.country_code
FROM events e1, events e2, venues v
WHERE e1.venue_id = v.venue_id
AND e2.venue_id = v.venue_id
AND DATE(e1.starts) = '2020-04-22'
AND DATE(e2.starts) = '2020-04-23';
```

```
SELECT v.name as Venue, v.country_code
FROM events e, venues v
WHERE e.venue_id = v.venue_id AND
DATE(e.starts) = '2020-04-22'
```

INTERSECT

```
SELECT v.name as Venue, v.country_code
FROM events e, venues v
WHERE e.venue_id = v.venue_id AND
DATE(e.starts) = '2020-04-23';
```

Aggregate Operators

```
SELECT count(title) FROM events
```

```
SELECT COUNT (DISTINCT title)
FROM events
```

```
SELECT min(starts), max(ends)
      FROM events INNER JOIN venues
      ON events.venue_id = venues.venue_id
      WHERE venues.name = 'Crystal Ballroom';
```

```
SELECT e.title FROM events e
      WHERE e.venue_id IN (SELECT v.venue_id FROM venues v
                           WHERE v.name = 'University of Iowa') AND
      (e.ends-e.starts) = (SELECT max(e2.ends-e2.starts)
                           FROM events e2 WHERE e2.venue_id IS NOT NULL);
```

```
COUNT (*)
COUNT ( [DISTINCT] A)
SUM ( [DISTINCT] A)
AVG ( [DISTINCT] A)
MAX (A) Can use Distinct
MIN (A) Can use Distinct
```

single column

Find date & title of the first event

- The first query is illegal! (wait for GROUP BY.)

```
SELECT e.title, min(e.starts)
FROM events e
```

How many tuples
in the result?

```
SELECT e.title, e.starts FROM events e
WHERE e.starts = (SELECT
min(e2.starts) FROM events e2);
```

GROUP BY and HAVING

- Apply aggregate to each of several *groups* of tuples
- Find the number of events registered *for each venue*

For *each venue i*:

```
SELECT COUNT (*)  
FROM events  
WHERE venue_id = i
```

```
SELECT venue_id, count(*)  
FROM events  
GROUP BY venue_id;
```

Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT] target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
HAVING      group-qualification
```

How many tuples
in the result?

- The **target-list** contains
 - Attribute names: must be a subset of *grouping-list*.
 - Terms with aggregate operations (e.g., COUNT (*)).
- The **group-qualification**
 - Must have a single value per group

Conceptual Evaluation

- Cross-product -> discard tuples -> apply projection
-> partition into groups using the *grouping-list* attribute values
-> eliminate groups that don't satisfy the *group-qualification*
- Expressions in *group-qualification* have a single value per group!
 - In effect, an attribute in *group-qualification* that is not an argument of an aggregate op also appears in *grouping-list*. (SQL does not exploit primary key semantics here!)
- One answer tuple is generated per qualifying group.


Find the venues with at least 2 events in 2019

```
SELECT venue_id, count(*)  
      FROM events  
     WHERE venue_id IS NOT NULL AND  
           EXTRACT(YEAR FROM starts)=2019  
     GROUP BY venue_id  
    HAVING count(*) >= 2;
```

Null Values

- Represent
 - *unknown* (e.g., rating not assigned) or
 - *inapplicable* (e.g., no spouse's name)
- Complications with nulls:
 - Operators to check if value is/is not *null*.
 - Is *rating* > 8 true or false when *rating* is null?
 - Answer: Evaluate to unknown
 - What about **AND**, **OR** and **NOT** connectives?
 - Need 3-valued logic (true, false and *unknown*)
 - Not unknown = unknown
 - WHERE clause eliminates rows that **don't evaluate to true**
 - New operators (in particular, *outer joins*) possible/needed.

p	q	p AND q	p OR q
T	T	T	T
T	F	F	T
T	U	U	T
F	T	F	T
F	F	F	F
F	U	F	U
U	T	U	T
U	F	F	U
U	U	U	U



Window functions – Partition by

- Not all RDBMS support it
- Similar to GROUP BY
 - Calculates aggregates on the OVER a PARTITION of the result set.
 - Rather than grouping the results outside of the SELECT attribute list, it returns grouped values as any other field

```
SELECT venue_id, count(*)  
      OVER (PARTITION BY venue_id)  
FROM events  
ORDER BY venue_id;
```

```
SELECT v.venue_id, v.name, e.title, count(*)  
      OVER (PARTITION BY v.venue_id)  
FROM events e INNER JOIN venues v  
ON e.venue_id = v.venue_id;
```