

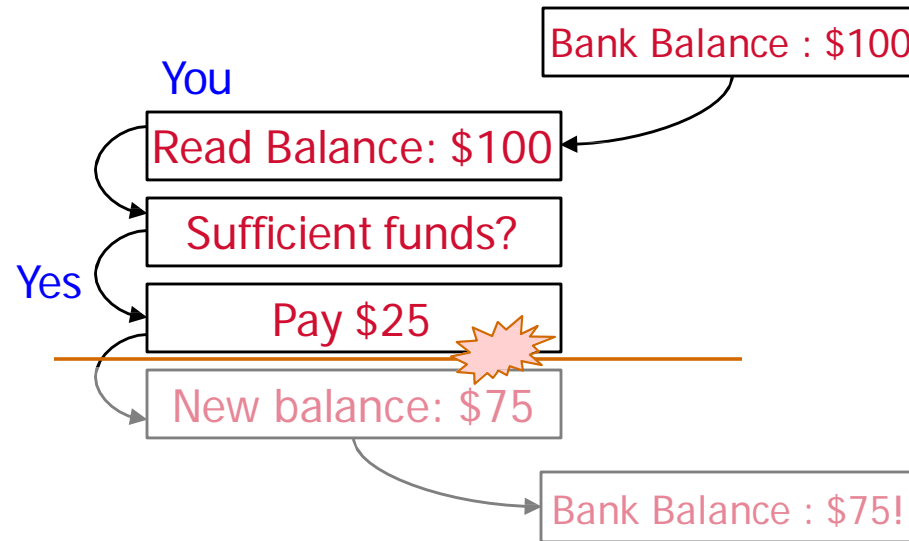
# **TRANSACTION MANAGEMENT**

# Transaction Management

```
Read (A);  
Check (A > $25);  
Pay ($25);  
A = A - 25;  

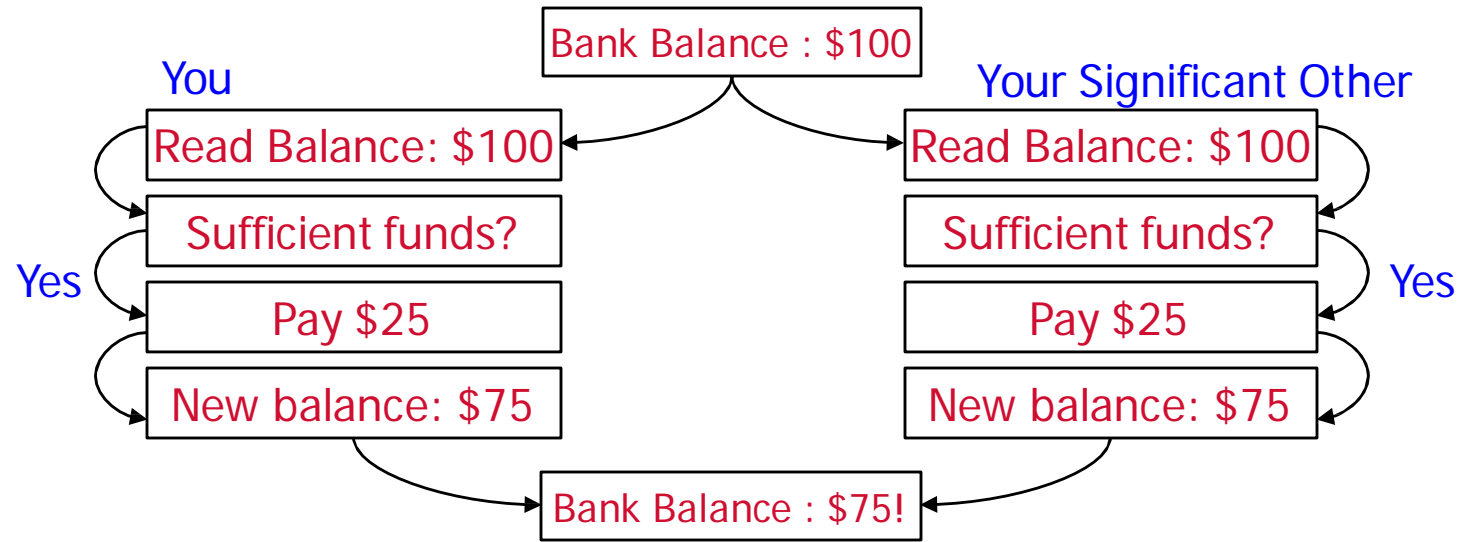

---

Write (A);
```



# Transaction Management

```
Read (A);  
Check (A > $25);  
Pay ($25);  
A = A - 25;  
Write (A);
```



- Inconsistency
  - Interleaving actions of different user programs
  - System crash/user abort/...
- Provide the users an illusion of a single-user system
  - Could insist on admitting only one query into the system at any time
    - lower utilization: CPU/IO overlap
    - long running queries starve other queries

# What is a Transaction?

- Collection of operations that form a single logical unit
  - A sequence of many actions considered to be one atomic unit of work
- Logical unit:
  - `begin transaction ... (SQL) end transaction`
- Operations:
  - Read (X), Write (X): Assume R/W on tuples (can be relaxed)
  - Special actions: `begin, commit, abort`
- Desirable Property: Must leave the DB in a consistent state
  - (DB is consistent when the transaction begins)
  - Consistency: DBMS only enforces integrity constraints (IC) specified by the user
  - DBMS does not understand any other semantics of the data

# The ACID Properties

TM

Xact. Mgmt.  
(logging)



**A**tomicity: All actions in the Xact happen, or none happen.

User



**C**onsistency: Consistent DB + consistent Xact  $\Rightarrow$  consistent DB

CC

Concurrency Ctrl.  
(locking)



**I**solation: Execution of one Xact is isolated from that of other Xacts.

RM

Recovery Mgmt.  
(WAL, ...)



**D**urability: If a Xact commits, its effects persist.

Begin

Read (A);

$A = A - 25;$

Write (A);

Read (B);

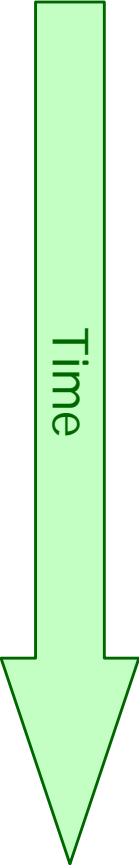
$B = B + 25;$

Write (B);

Commit

# Schedules

- Schedule: An interleaving of actions from a set of Xacts, where the actions of any one Xact are in the original order.
  - Actions of Xacts *as seen by the DB*
  - *Complete schedule* : each Xact ends in commit or abort
  - *Serial schedule* : No interleaving of actions from different Xacts.
- Initial State + Schedule → Final State



I1	I2
begin	
R(A)	
W(A)	
	begin
	R(B)
	W(B)
R(C)	
W(C)	
	commit
abort	

# Acceptable Schedules

- One sensible “isolated, consistent” schedule:
  - Run Xacts one at a time (serial schedule)
- Serializable schedules:
  - Final state is what ***some complete*** serial schedule of ***committed*** transactions would have produced.
  - Can different serial schedules have different final states?
    - Yes, all are “OK”!
  - Aborted Xacts?
    - ignore them for a little while (made to ‘disappear’ using logging)
  - Other external actions (besides R/W to DB)
    - e.g. print a computed value, fire a missile, ...
    - Assume (for this class) these values are written to the DB, and can be undone

# Serializability Violations

- @Start (A,B) = (1000, 100)
  - End (990, 210)
- T1 → T2:
  - (900, 200) → (990, 220)
- T2 → T1:
  - (1100, 110) → (1000, 210)
- **W-R conflict:** Dirty read
  - *Could* lead to a non-serializable execution
- Also R-W and W-W conflicts

<i>T1: Transfer \$100 from A to B</i>	<i>T2: Add 10% interest to A &amp; B</i>
begin	
	begin
R(A) /A -= 100	
<b>W(A)</b>	
	<b>R(A)</b> /A *= 1.1
	W(A)
	R(B) /B *= 1.1
	W(B)
	commit
R(B) /B += 100	
W(B)	
commit	

Database  
Inconsistent



# More Conflicts

- RW Conflicts (Unrepeatable Read)
  - $R_{T2}(X) \rightarrow W_{T1}(X)$ , T1 overwrites what T2 read.
  - $R_{T2}(X) \rightarrow W_{T1}(X) \rightarrow R_{T2}(X)$ . T2 sees a different X value!
- WW Conflicts (Overwriting Uncommitted Data)
  - T2 overwrites what T1 wrote.
    - E.g. : Students in the same group get the same project grade.
    - $T_p: W(X=A), W(Y=A)$        $T_{TA}: W(X=B), W(Y=B)$
    - $W_p(X=A) \rightarrow W_{TA}(X=B) \rightarrow W_{TA}(Y=B) \rightarrow W_p(Y=A)$   
[Note: no reads]
  - Usually occurs in conjunction with other anomalies.

# Now, Aborted Transactions

- Serializable schedule: Equivalent to a serial schedule of *committed* Xacts.
  - as if aborted Xacts *never happened*.
- Two Issues:
  - How does one undo the effects of a Xact?
    - Logging/recovery
  - What if another Xact sees these effects??
    - Must undo that Xact as well!

# Cascading Aborts

- Abort of T1 requires abort of T2!
  - Cascading Abort

T1	T2
begin	
R(A)	
W(A)	
	begin
	R(A)
	W(A)
	commit
abort	

# Cascading Aborts

- Abort of T1 requires abort of T2!
  - **Cascading Abort**
- Consider commit of T2
  - Can we undo T2?
- *Recoverable* schedule: Commit only after all xacts that supply dirty data have committed.

T1	T2
begin	
R(A)	
W(A)	
	begin
	R(A)
	W(A)
commit	
	commit

# Cascading Aborts

- *ACA (avoids cascading abort)* schedule
  - Transaction only reads committed data
  - One in which cascading abort cannot arise.
  - Schedule is also recoverable

<u>T1</u>	<u>T2</u>
begin	
R(A)	
W(A)	
commit	
	begin
	R(A)
	W(A)
	Commit

<u>T1</u>	<u>T2</u>
begin	
R(A)	
W(A)	
	begin
	R(A)
	W(A)
abort	
	commit

# Locking: A Technique for C. C.

- Concurrency control usually done via locking.
- Lock info maintained by a “lock manager”:
  - Stores (XID, RID, Mode) triples.
  - This is a simplistic view; suffices for now.
  - Mode  $\in \{S, X\}$
  - Lock compatibility table:
- If a Xact can't get a lock
  - Suspended on a wait queue
- When are locks acquired?
  - Buffer manager call!

	--	S	X
--	✓	✓	✓
S	✓	✓	
X	✓		


# Two-Phase Locking (2PL)

- **2PL:**

- If T wants to read (modify) an object, first obtains an S (X) lock
- If T releases any lock, it can acquire no new locks!
- **Guarantees serializability! Why?**

- **Strict 2PL:**

- **Hold all locks until end of Xact**
- Guarantees serializability, and ACA too!
  - Note ACA schedules are always recoverable



# Schedule with Locks


<i>T1: Transfer \$100 from A to B</i>	<i>T2: Add 10% interest to A &amp; B</i>
begin	
	begin
R(A) /A -= 100	
<b>W(A)</b>	
	<b>R(A)</b> /A *= 1.1
	W(A)
	R(B) /B *= 1.1
	W(B)
	commit
R(B) /B += 100	
W(B)	
commit	

<i>T1</i>	<i>T2</i>
begin	
	begin
<b>X(A)</b>	
R(A)	
<b>W(A)</b>	
	<b>X(A) – Wait!</b>
<b>X(B)</b>	
R(B)	
W(B)	
U <sub>x</sub> (A), U <sub>x</sub> (B)/commit	
	<b>R(A)</b>
	W(A)
	...



# Deadlocks

$X_{T1}(B), X_{T2}(A), S_{T1}(A), S_{T2}(B)$



- Deadlocks can cause the system to wait forever.
- Need to detect deadlock and break, or prevent deadlocks
- Simple mechanism: timeout and abort
- More sophisticated methods exist

# Conflict Serializability & Graphs


Theorem: A schedule is conflict serializable iff its precedence graph is acyclic

Theorem: 2PL ensures that the precedence graph will be acyclic

- Why Strict 2PL?
  - Guarantees ACA
    - read only committed values
  - How? Write locks until EOT
    - No WW or WR => on abort replace original value

# Deadlocks


$X_{T1}(B), X_{T2}(A), S_{T1}(A), S_{T2}(B)$



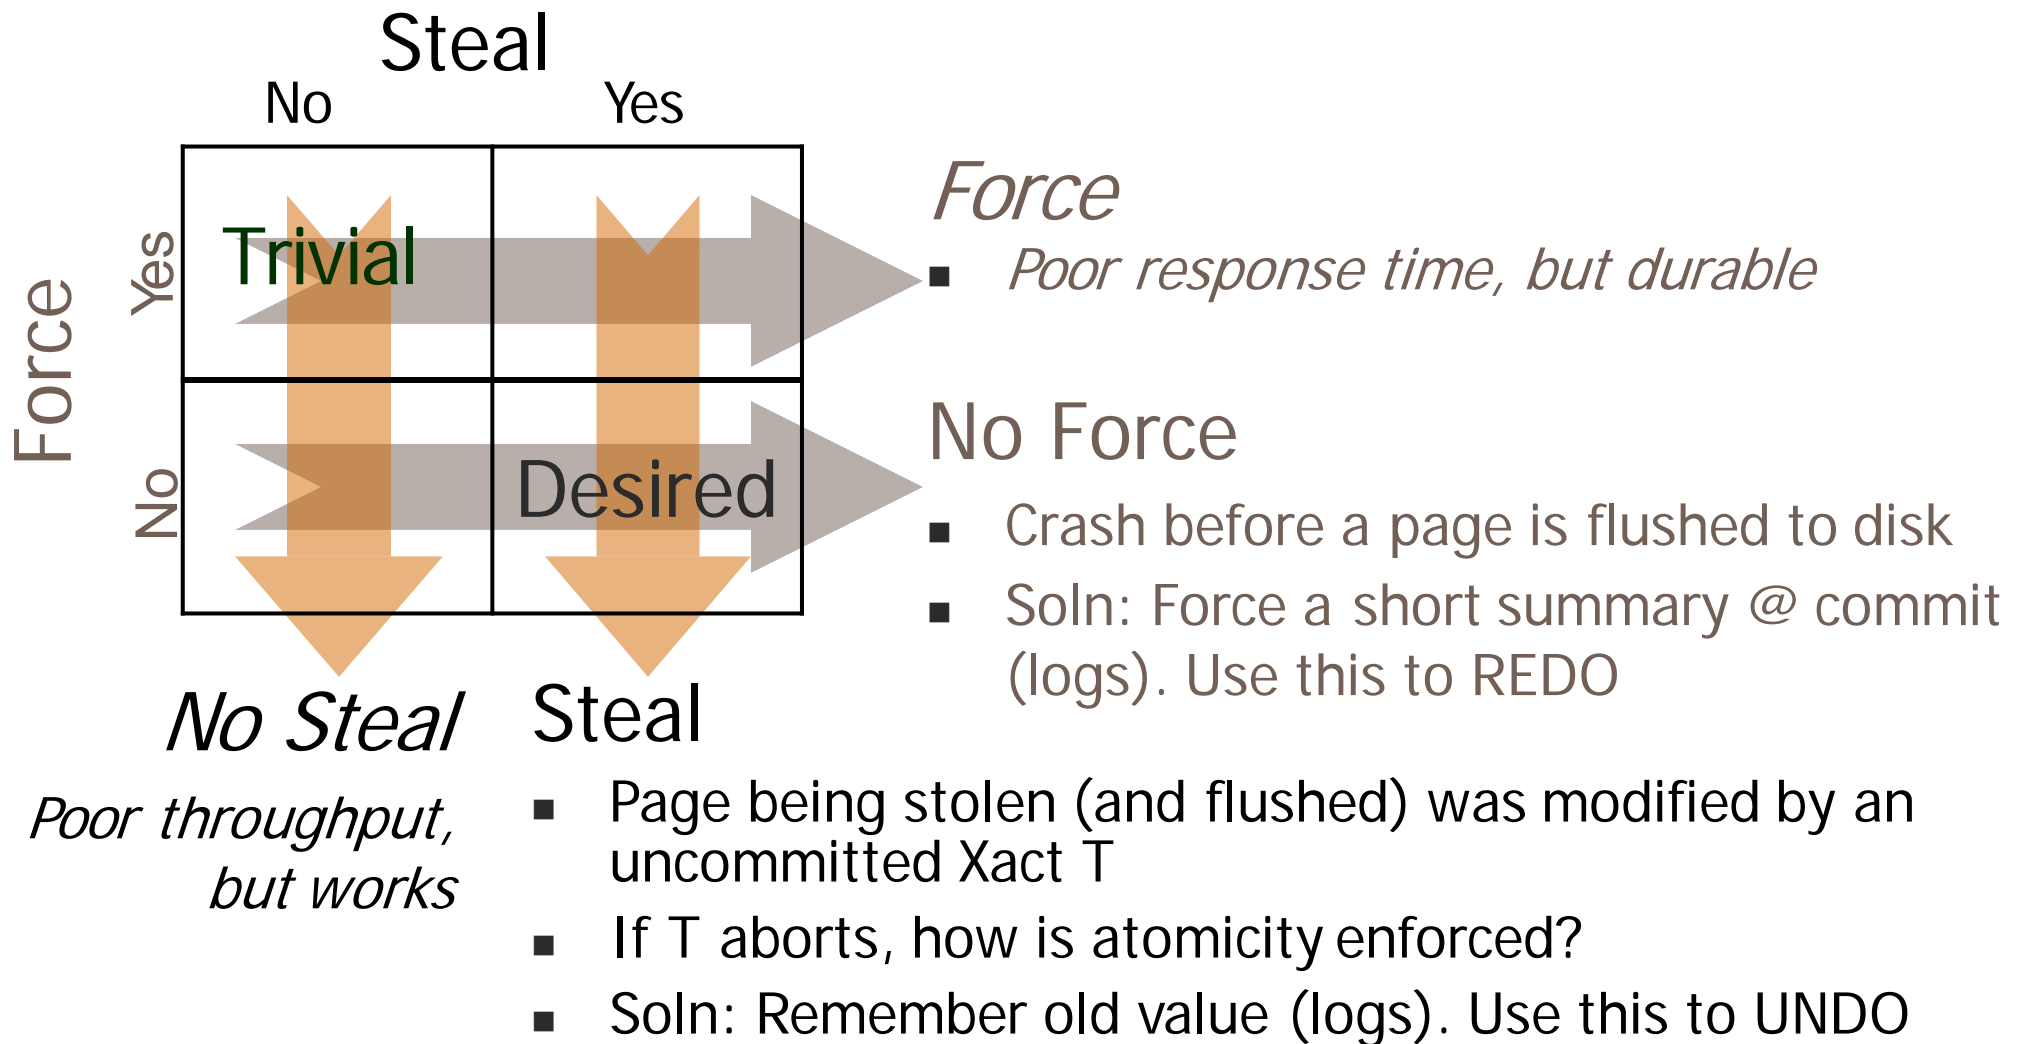
- Deadlocks can cause the system to wait forever.
- Need to detect deadlock and break, or prevent deadlocks
- Detect deadlock
  - Draw a lock graph. Cycles implies a deadlock
- Alternative ways of dealing with deadlock
  - Break Deadlock
    - On each lock request “update the lock graph”. If a cycle is detected, abort one of the transactions. The aborted transaction is restarted after waiting for a time-out interval.
  - Prevent deadlock
    - Assign priorities to the transactions. If a transaction, T1, requests a lock that is being held by another transaction, T2, with a lower priority, then T1 “snatches” the lock from T2 by aborting T2 (which frees up the lock on the resource). T2 is then restarted again after a time-out.

# Ensuring Atomicity & Durability

- Atomicity:
  - Transactions may abort (“Rollback”) -> no effects should be seen
- Durability:
  - What if DBMS stops running? (Causes?)
- Desired Behavior after system restarts:
  - T1, T2 & T3 should be durable.
  - T4 & T5 should be aborted



# Buffer Pool: Sharing & Writing



# Basic Idea: Logging

- Record information, for every change, in a *log*.
  - Sequential writes to log (put it on a separate disk).
  - Stored in stable storage to survive system crash
    - disk mirroring
  - Each record has a log sequence number (LSN)
  - Log record contains:
    - <prevLSN, XID, type, ... >
    - and additional control info (which we'll see soon)
    - Note: the log records for a transaction are chained by prevLSN

# Write-Ahead Logging (WAL)

- The **Write-Ahead Logging** Protocol:
  1. Must force the log record for an update before the corresponding data page gets to storage.
  2. Must write all log records for a Xact before commit.
- #1 guarantees Atomicity.
- #2 guarantees Durability.

If DB says TX **commits**, TX effect **remains** after database crash

DB can **undo actions** and help us with **atomicity**

# Normal Execution of a Xact

- Series of **reads & writes**, followed by **commit** or **abort**.
  - Updates are “in place”: i.e., data on disk is overwritten
  - We will assume that write is atomic on disk.
    - In practice, additional details to deal with non-atomic writes.
- **Strict 2PL.**
- STEAL, NO-FORCE buffer management, with **Write-Ahead Logging.**



# The ACID Properties

TM

Xact. Mgmt.  
(logging)



**A**tomicity: All actions in the Xact happen, or none happen.

User



**C**onsistency: Consistent DB + consistent Xact  $\Rightarrow$  consistent DB

CC

Concurrency Ctrl.  
(locking)



**I**solation: Execution of one Xact is isolated from that of other Xacts.

RM

Recovery Mgmt.  
(WAL, ...)



**D**urability: If a Xact commits, its effects persist.

# Postgres

- PostgreSQL transactions follow ACID compliance
- Atomicity & Durability: WAL logging, Continuous Archiving and Point-in-Time Recovery (PITR)
- Consistency: Multiversion Concurrency Control, MVCC
- Isolation: *Serializable Snapshot* Isolation (SSI)
- Every command we've executed in psql has been implicitly wrapped in a transaction
- Explicit transaction:  
    **BEGIN TRANSACTION;**  
    **DELETE FROM events;**  
    **ROLLBACK;**  
    **SELECT \* FROM events;**  
  
    **END;** or **COMMIT;** Commits the current transaction

# Install MongoDB

We will be using the free MongoDB Community Edition Database server.

<https://www.mongodb.com/try/download/community>

Download a build for your OS

After download completes, follow the Installation instructions

<https://docs.mongodb.com/manual/installation/>

From a terminal window in your computer, run the MongoDB server (`mongod`).

Then run the mongo shell client (`mongo`) -

<https://docs.mongodb.com/manual/mongo/>

Take a screenshot of your terminal window and upload to ICON in the dropbox.

To exit the shell, type **quit()** or use the **<Ctrl-C>** shortcut