

Key-value stores

Redis

<http://redis.io>

The Key-value Abstraction

- (Business) Key \rightarrow Value
- (twitter.com) tweet id \rightarrow information about tweet
- (amazon.com) item number \rightarrow information about it
- (kayak.com) Flight number \rightarrow information about flight, e.g., availability
- (yourbank.com) Account number \rightarrow information about it

The Key-value Abstraction (2)

- It's a dictionary data structure.
 - Insert, lookup, and delete by key
 - E.g., hash table, binary tree
- But distributed.
- Sound familiar?
 - Distributed Hash tables (DHT) in P2P systems
- It's not surprising that key-value stores reuse many techniques from DHTs.

Key Value Stores

- Key-Valued data model
 - Key is the unique identifier
 - Key is the granularity for consistent access
 - Value can be structured or unstructured
- Gained widespread popularity
 - In house: **Bigtable** (Google), **PNUTS** (Yahoo!), **Dynamo** (Amazon)
 - Open source: **Redis**, **HBase**, **Hypertable**, **Cassandra**, **Voldemort**
- Popular choice for the modern breed of web-applications

Important Design Goals

- **Scale out: designed for scale**
 - Commodity hardware
 - Low latency updates
 - Sustain high update/insert throughput
- **Elasticity – scale up and down with load**
- **High availability – downtime implies lost revenue**
 - Replication (with multi-mastering)
 - Geographic replication
 - Automated failure recovery

Lower Priorities

- **No Complex querying functionality**
 - No support for SQL
 - CRUD operations through database specific API
- **No support for joins**
 - Materialize simple join results in the relevant row
 - Give up normalization of data?
- **No support for transactions**
 - Most data stores support single row transactions
 - Tunable consistency and availability
- ***Avoid scalability bottlenecks at large scale***

System Interface

- Two basic operations:
 - Get(key):
 - Put(key, value)

Redis



- Redis is an open source, advanced **key-value data store**
- Often referred to as a **data structure server** since keys can contain strings, hashes, lists, sets and sorted sets
- The name Redis means Remote Dictionary Server
- Redis works with an **in-memory** dataset
- It is possible to **persist** dataset either by
 - dumping the dataset to disk every once in a while
 - or by appending each command to a log

Who is using Redis?

- [Twitter](#)
- [GitHub](#)
- [Weibo](#)

- [Pinterest](#)
- [Snapchat](#)
- [Craigslist](#)

- [Digg](#)
- [StackOverflow](#)
- [Flickr](#)



Configuration

- Configuration file: `/redis/redis.conf`
- It is possible to change a port (if you wish):

```
port 6379
```

- For development environment it is useful to change data persisting policy

```
save 900 1  
save 300 10  
save 60 10000
```



```
save 10 1
```


save after 10 sec if at least 1 key changed

Running Redis Server

- Run `/redis/bin/redis-server.exe` and specify configuration file to use

```
redis>redis-server redis.conf
```

```
C:\tmp\redis>redis-server redis.conf
```



```
Redis 2.6.12 (00000000/0) 64 bit
```

```
Running in stand alone mode
```

```
Port: 6379
```

```
PID: 10720
```

```
http://redis.io
```

```
[10720] 24 Oct 11:49:43.565 # Server started, Redis version 2.6.12
```

```
[10720] 24 Oct 11:49:43.565 * The server is now ready to accept connections on port 6379
```

Running Redis Client

- Run `redis-cli.exe`
- Now we can play with Redis a little bit

```
C:\tmp\redis>redis-cli
redis 127.0.0.1:6379> SET MyVar 10
OK
redis 127.0.0.1:6379> GET MyVar
"10"
redis 127.0.0.1:6379> INCR MyVar
(integer) 11
redis 127.0.0.1:6379> INCRBY MyVar 10
(integer) 21
```

Useful Commands

- Print all keys:

```
KEYS *
```

- Remove all keys from all databases

```
FLUSHALL
```

- Synchronously save the dataset to disk

```
SAVE
```

Redis keys

- Keys are binary safe - it is possible to use any binary sequence as a key
- The empty string is also a valid key
- Too long keys are not a good idea
- Too short keys are often also not a good idea ("`u:1000:pwd`" versus "`user:1000:password`")
- Nice idea is to use some kind of schema, like: "`object-type:id:field`"

Redis data types

Redis is often referred to as a **data structure server** since keys can contain:

- Strings
- Lists
- Sets
- Hashes
- Sorted Sets

Redis Strings

- Most basic kind of Redis value
- Binary safe - can contain any kind of data, for instance a JPEG image or a serialized Ruby object
- Max 512 Megabytes in length
- Can be used as atomic counters using commands in the INCR family
- Can be appended with the APPEND command

Redis Strings: Example

```
redis 127.0.0.1:6379> SET COUNTER 10
OK
redis 127.0.0.1:6379> INCRBY COUNTER 100
(integer) 110
redis 127.0.0.1:6379> DECR COUNTER
(integer) 109
redis 127.0.0.1:6379> APPEND COUNTER 01
(integer) 5
redis 127.0.0.1:6379> GET COUNTER
"10901"
redis 127.0.0.1:6379> INCR COUNTER
(integer) 10902
```

Transactions

- Redis's MULTI block atomic commands are a similar concept to transactions. Wrapping two operations like SET and INCR in a single block will complete either successfully or not at all.
- We begin the transaction with the MULTI command and execute it with EXEC (rollback with DISCARD).

```
redis 127.0.0.1:6379> MULTI
redis 127.0.0.1:6379> SET foo bar
redis 127.0.0.1:6379> INCR counter
redis 127.0.0.1:6379> EXEC
```

Redis Hashes

- Redis objects that can take any number of key-value pairs
- Map between string fields and string values
- Perfect data type to represent objects

```
HMSET user:1000 username gomez password P1pp0 age 34
HGETALL user:1000
HVALS user:1000
HKEYS user:1000
HSET user:1000 password 12345
HGETALL user:1000
HGET user:1000 username
```

Redis Lists

- Lists of ordered values (insertion order)
- Can act as queues or stacks (or just lists)
- Add elements to a Redis List pushing new elements on the head (on the left) or on the tail (on the right) of the list
- Max length: $(2^{32} - 1)$ elements
- Model a timeline in a social network, using LPUSH to add new elements, and using LRANGE in order to retrieve recent items
- Use LPUSH together with LTRIM to create a list that never exceeds a given number of elements

Redis Lists: Example

```
redis 127.0.0.1:6379> LPUSH myList a
(integer) 1
redis 127.0.0.1:6379> LPUSH myList b
(integer) 2
redis 127.0.0.1:6379> LPUSH myList c
(integer) 3
redis 127.0.0.1:6379> LLEN myList
(integer) 3
redis 127.0.0.1:6379> LRANGE myList 0 -1
1) "c"
2) "b"
3) "a"
redis 127.0.0.1:6379> RPUSH myList d e f
(integer) 6
redis 127.0.0.1:6379> LRANGE myList 0 -1
1) "c"
2) "b"
3) "a"
4) "d"
5) "e"
6) "f"
redis 127.0.0.1:6379> LTRIM myList 2 4
OK
redis 127.0.0.1:6379> LRANGE myList 0 -1
1) "a"
2) "d"
3) "e"
```

More list functions

- LREM removes from the list given value
LREM myList 0 a
- LPOP removes from the left (head) of the list
LPOP myList
- RPUSH/RPOP add/remove from the right of the list
- RPOPLPUSH pop a value from the tail of one list and push it to the head of another
RPOPLPUSH myList yourList

Blocking lists

- Producer-consumer example.
- Open another redis client, one client (the consumer) just listens for new comments and pop them as they arrive.
BRPOP comments 300
- The command will block until a value exists to pop. Timeout in seconds is set to five minutes.
- Now the producer should push a message to comments.
LPUSH comments "ModernDB is a great class!"
- Switch back to the consumer console, two lines will be returned: the key and the popped value. The console will also output the length of time it spent blocking.

Sets

- Unordered collections with no duplicate values, supports unions and intersections

```
SADD myPref movies reading walking
```

- SMEMBERS retrieves the whole set.

```
SMEMBERS myPref
```

```
SADD yourPref running painting reading fishing
```

- To find the intersection use the SINTER command.

```
SINTER myPref yourPref
```

- Remove any matching values in one set from another:

```
SDIFF myPref yourPref
```

- Union is a set, any duplicates are dropped.

```
SUNION myPref yourPref
```

- That set of values can also be stored directly into a new set:

```
SUNIONSTORE hobbies myPref yourPref
```


Redis Sorted Sets

- Every member of a Sorted Set is associated with score, that is used in order to take the sorted set ordered, from the smallest to the greatest score
`ZADD scoreboard 500 me 9 you 15 him`
- To increment a score, we can either re-add it with the new score, which just updates the score but does not add a new value, or increment by some number, which will return the new value.

```
ZINCRBY scoreboard 1 you
```

Sorted Set Ranges

- To get scores from the sorted set:
`ZRANGE scoreboard 0 1`
- To get scores append `WITHSCORES`
- `ZREVRANGE` gets them in reverse
- `ZRANGEBYSCORE` allow to provide score range (inclusive by default)

```
ZRANGEBYSCORE scoreboard 100 500
```

```
ZRANGEBYSCORE scoreboard (100 500
```

```
ZRANGEBYSCORE scoreboard (100 inf
```

Sorted Set Unions

```
ZUNIONSTORE destination numkeys key [key ...]  
    [WEIGHTS weight [weight ...]] [AGGREGATE SUM|MIN|MAX]
```

- destination is the key to store into
- numkeys is simply the number of keys you're about to join
- key is one or more keys to union
- weight [optional] is the number to multiply each score of the relative key by (if you have two keys, you can have two weights, and so on).
- aggregate is the optional rule, sum is default

For today

- Add a sorted set called scoreboard with the values 50 me, 30 you, 15 her, and 10 him
- Add the gamesWon sorted set for three players: you (3), me (4), and him (8)
- Multiply by 10 the number of wins as the points for each player and add them to the scoreboard (can be done in one or more commands)
- Retrieve the scoreboard in reverse order
- Submit your commands and outputs to ICON