

# Advanced usage and Distribution Redis

<http://redis.io>

# Expiry


- EXPIRE command, an existing key, and a time to live (in seconds).
- Set ice to expire after 10 seconds
- SET ice "I'm melting..."
- EXPIRE ice 10
- EXISTS ice
- Wait 10 secs
- EXISTS ice
- Shortcut command:
- SETEX ice 10 "I'm melting..."
- Check time to live:
- TTL ice
- Make it persistent
- PERSIST ice

# Database Namespaces

- In Redis a namespace is called a *database* and is keyed by number
- So far we've always interacted with the default database (namespace 0)
- SET greeting hello
- GET greeting
- SELECT 1
- GET greeting
- SET greeting hola
- SELECT 0
- GET greeting
- Since all databases are running in the same server instance, you can shuffle keys around using MOVE
- MOVE greeting 2
- SELECT 2
- GET greeting

# Publish-subscribe

- Multiple subscribers
- Start two clients:
- SUBSCRIBE comments
- Start publisher:
- PUBLISH comment “Thanksgiving is next week!”
- UNSUBSCRIBE disconnects client – in redis-cli console press Ctrl+C to break the connection



# Redis configuration

- Server info: INFO
- daemonize no – starts in the foreground
- port 6379
- loglevel verbose | notice | warning
- logfile stdout /\*filename req if daemonize mode\*/
- database 16

# Redis configuration (cont.)

- `save 300 1` (snapshotting, save every 5 mins if any keys change at all)
- `appendonly yes` (keeps record of all write commands)
  - `appendfsync` `always` | `everysec` | `no`

# Security

- Redis is not natively built to be a fully secure server
- Plaintext password not really safe
- Use SSH security
- Redis allows you hide or suppress commands
  - Include rename-command in the conf file:
  - rename-command FLUSHALL c123456789
  - rename-command FLUSHALL ""

# Master-slave replication

- First make a copy of the redis.conf file
  - `cp redis.conf redis-s1.conf`
  - Change port and slaveof
    - port 6380
    - Slaveof 127.0.0.1 6379
  - Start both servers
    - `Redis-server redis-s1.conf`
  - Add to the server
    - `SADD meetings "Initial group meeting" "ECE potluck"`
  - Query in the slave
    - `SMEMBERS meetings`



# Redis cluster

- Many Redis clients provide an interface for building a simple ad hoc distributed Redis cluster.
- Unlike the master-slave setup, both of servers take the master (default) configuration.


```
redis = Redis::Distributed.new([  
  "redis://localhost:6379/" ,  
  "redis://localhost:6380/"  ])
```

- Consistent hashing is used to manage the cluster

# Consistent Hashing System

- Given  $k$ , every node can locate  $n_k$
- Hash every node's IP address
  - map these values on a circle
- Given a key  $k$ , hash  $k$ 
  - $k$  is assigned to closest node on circle, moving clockwise.

# Consistent Hashing System




# Consistent Hashing System

- Pros:
  - Load Balanced
  - Dynamic membership
    - when  $N^{\text{th}}$  node joins network, only  $O(1/N)$  keys are moved to rebalance
- Con:
  - Every node must know about every other node
  - $O(N)$  memory,  $O(1)$  communication
    - Not scalable in number of nodes


# Scaling Consistent Hashing

- Approach 0:
  - Each node keeps track of only their successor
  - Resolution of hash function done through routing
  - $O(1)$  memory
  - $O(N)$  communication
- Approach 1:
  - Each node keeps track of  $O(\log N)$  successors in a “finger table”
  - $O(\log N)$  memory
  - $O(\log N)$  communication

# Finger Table Pointers



# Routing with Finger Tables



# Incremental Scalability

- Utilize “virtual nodes” along ring
  - Many virtual nodes per physical node
  - larger machines can hold more virtual nodes
  - Heterogeneous hardware is properly load balanced

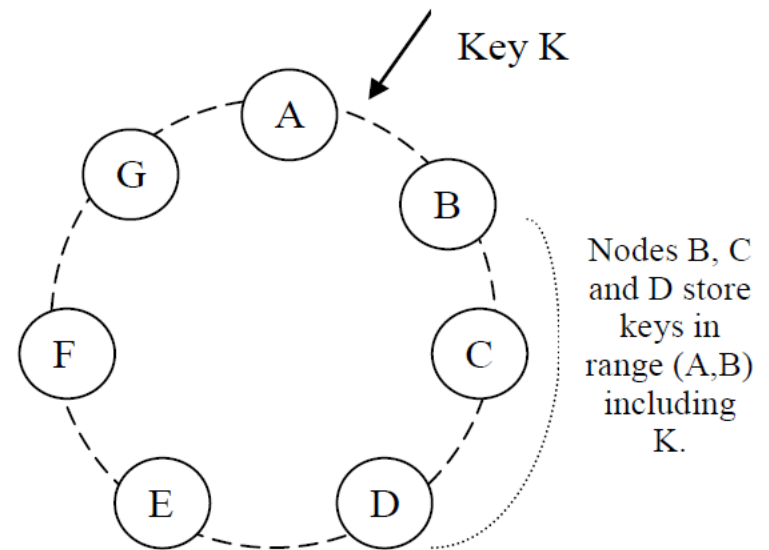


# Replication

Each data item is **replicated** at N hosts.

*preference list*: The list of **nodes that is responsible for storing** a particular key.

Some fine-tuning to account for virtual nodes



# Preference Lists

- List of nodes **responsible for storing** a particular key.
- Due to **failures**, preference list contains **more than N nodes**.
- Due to **virtual** nodes, preference list skips positions to ensure **distinct physical nodes**.

# Replication: Sloppy Quorum

- **Quorum System:**  $R + W > N$ ,  $W > N/2$ 
  - $R$ ,  $W$ ,  $N$  are tunable
- Each node maintains a “preference list” of replicas for its own data
- Replicas are made on first  $N$  *healthy* nodes from preference list
  - require  $R$  nodes to respond for `get()`
  - require  $W$  nodes to respond for `put()`

# Data Versioning

- A `put()` call may **return to its caller before** the update has been applied at all the replicas
- A `get()` call may **return many versions** of the same object.
- **Challenge:** an object may have distinct versions
- **Solution:** use **vector clocks** in order to capture **causality** between different versions of same object.

# Vector Clock

- A **vector clock** is a list of (node, counter) pairs.
- **Every version** of every object is associated with **one** vector clock.
- If the **all counters** on the first object's clock are **less-than-or-equal** to **all** of the **counters** in the **second clock**, then the **first is an ancestor of the second** and can be **forgotten**.
- **Application reconciles divergent versions** and collapses into a single new version.

# High Availability for Writes

- Clients write to first node they find
  - Vector clocks timestamp writes
  - Different versions of key's value live on different nodes
- Conflicts are resolved during reads
  - Like git: “automerge conflict” is handled by end application

# Using Redis from your program

- Java Clients for Redis
  - **Lettuce** or **Jedis**
    - <https://redislabs.com/lp/redis-java/>
- Python Client for Redis
  - **Redis-py**
    - <https://redislabs.com/lp/python-redis/>

# For today...

- Install your favorite programming language driver and connect to the Redis server.
- Insert and increment a value within a transaction.
- Submit the file with your code to ICON.